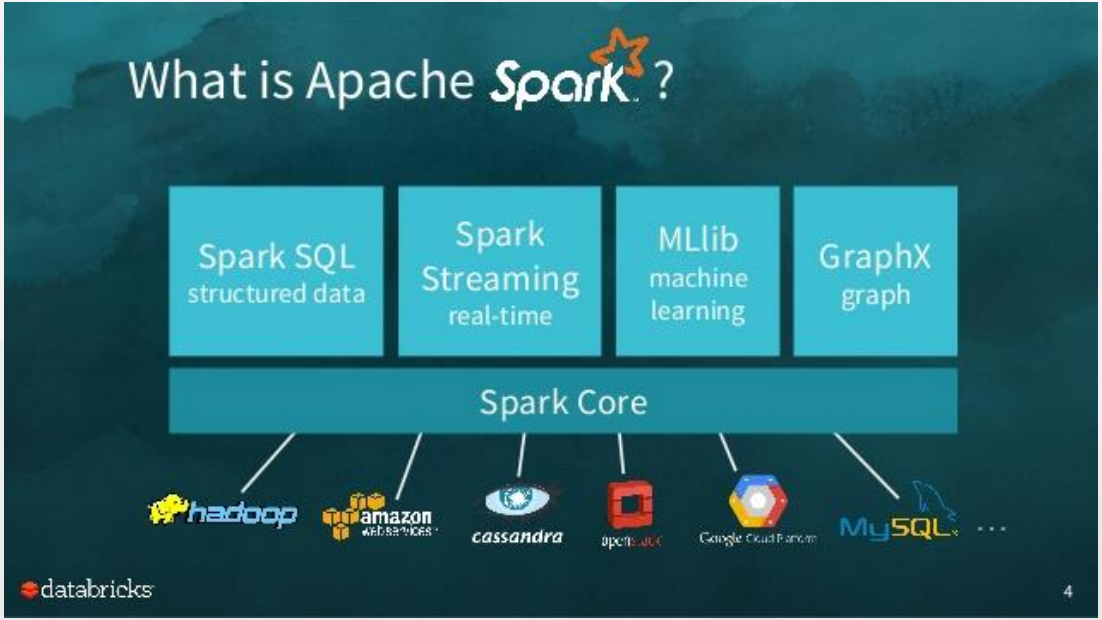

The Pilot Of Spark

2017.5 XenRon

http://spark.apache.org/docs/latest/

CONTENTS

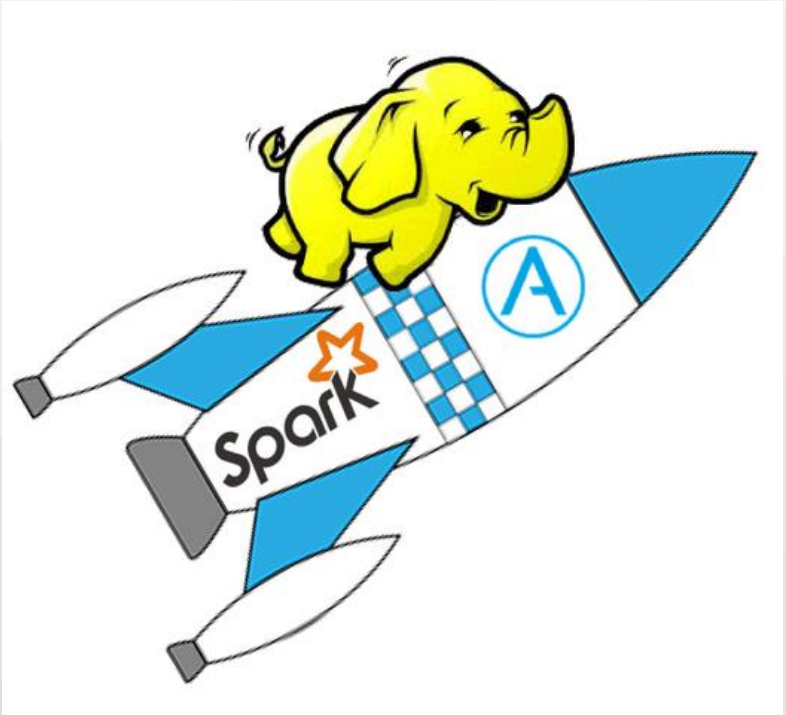

Review

Map Reduce

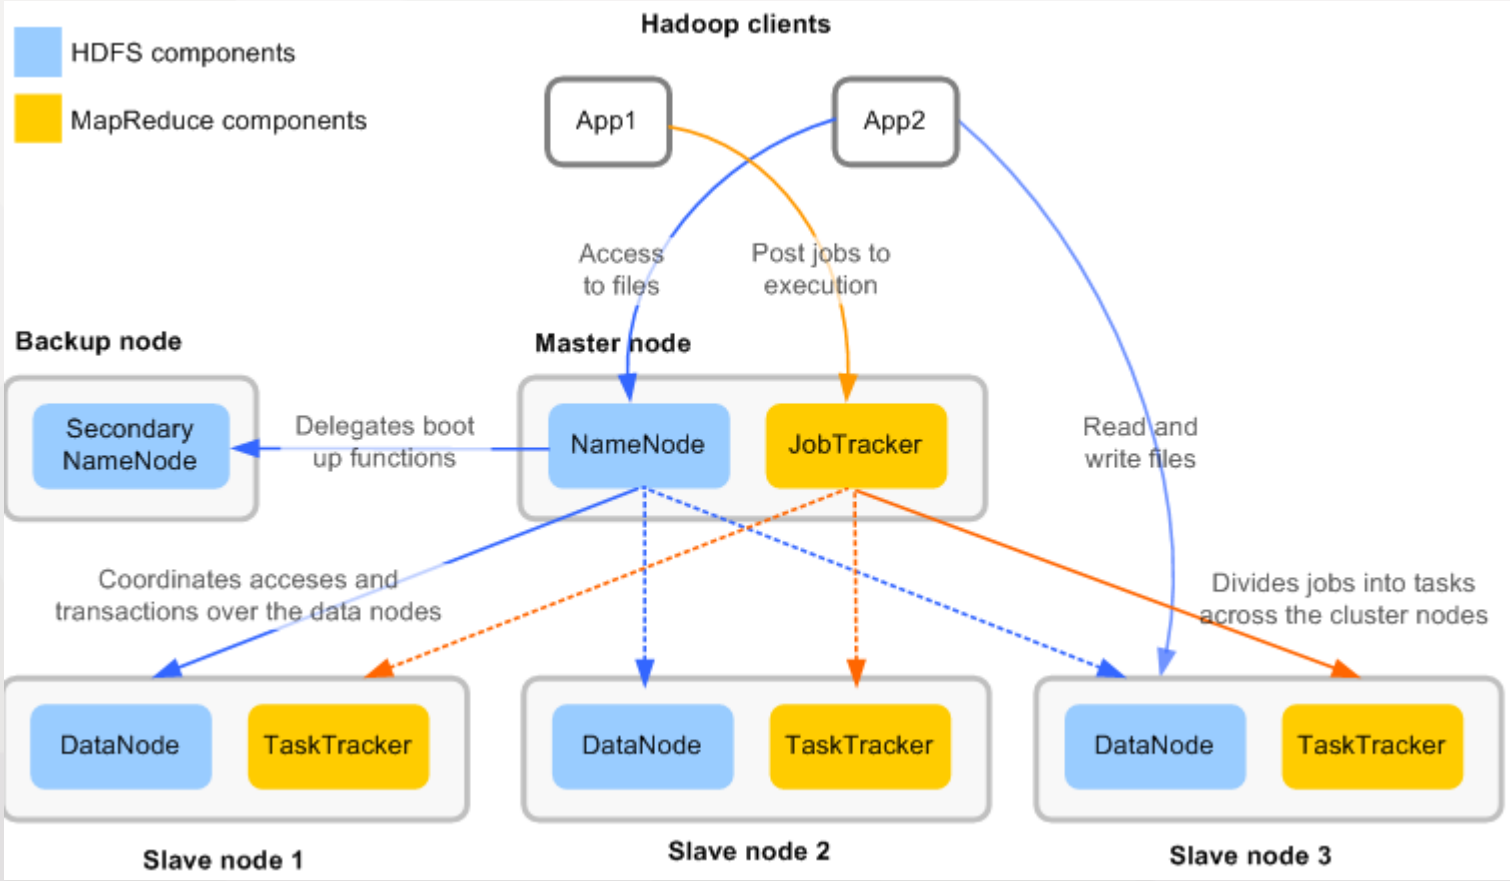
Hadoop clients
HDFS components
MapReduce components
App1
App2
Access to files
Post jobs to execution
Backup node
Master node
Secondary NameNode
Delegates boot up functions
NameNode
JobTracker
Read and write files
Coordinates acceses and transactions over the data nodes
Divides jobs into tasks across the cluster nodes
DataNode
TaskTracker
DataNode
TaskTracker
DataNode
TaskTracker
Slave node 1
Slave node 2
Slave node 3

HDFS

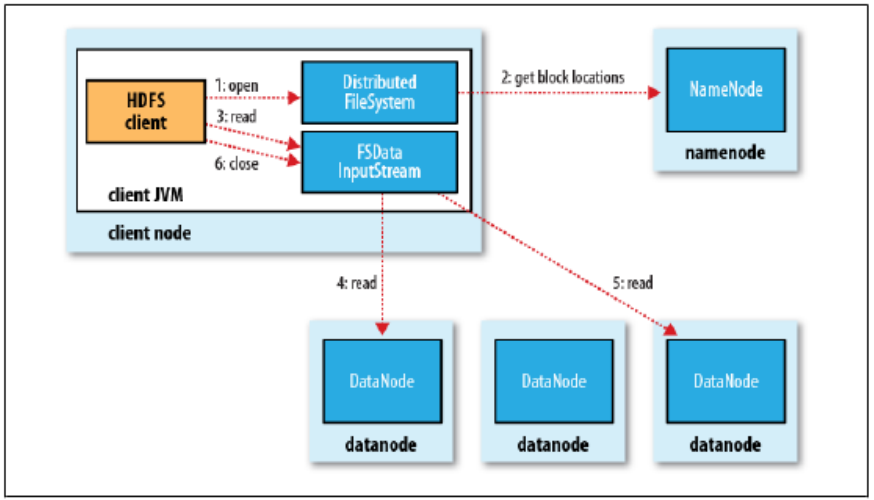

Figure 3-2. A client reading data from HDFS

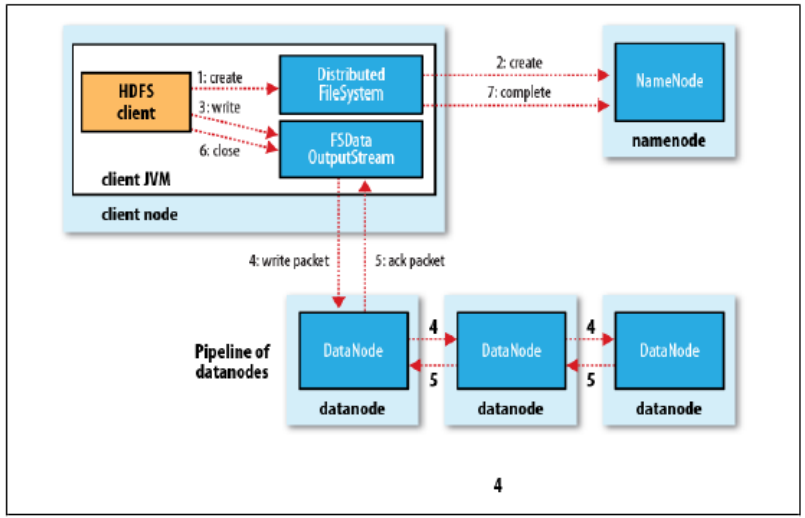

Figure 3-4. A client writing data to HDFS

Map Reduce

Hadoop 1.x

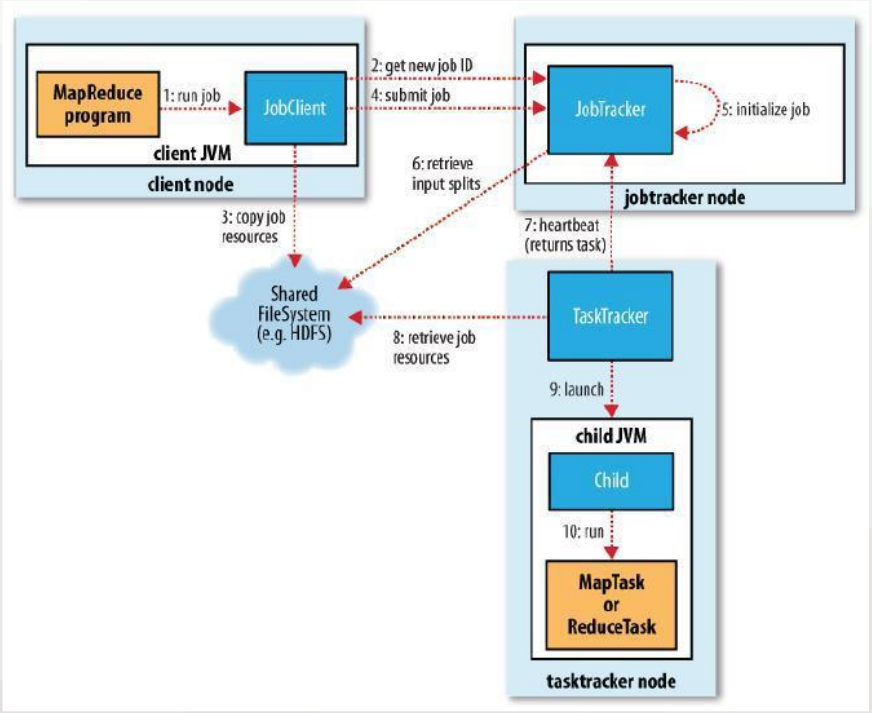

Hadoop 2.x

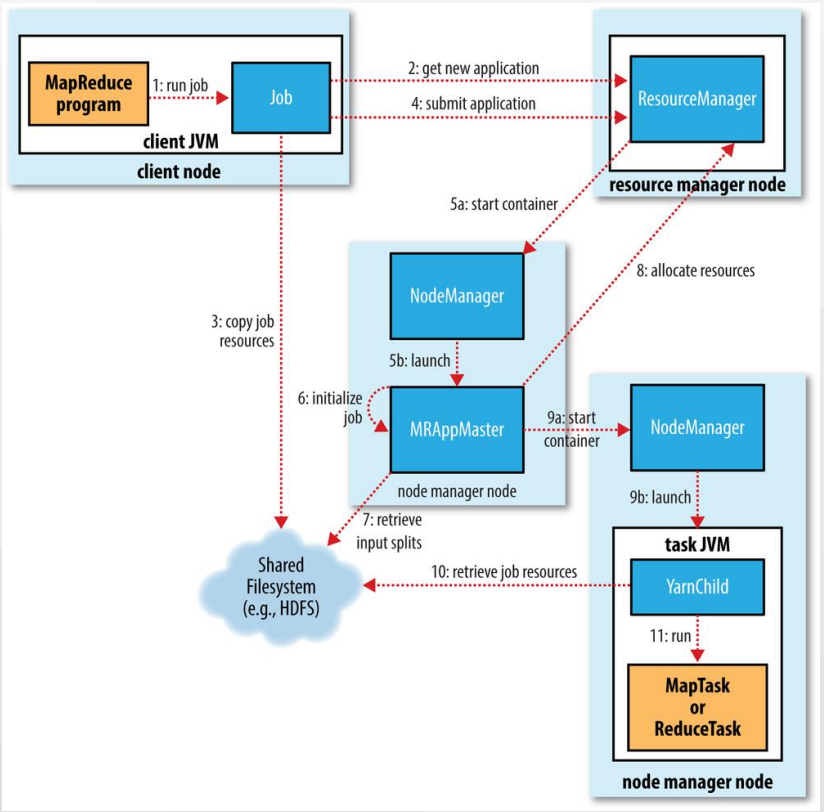

HBase

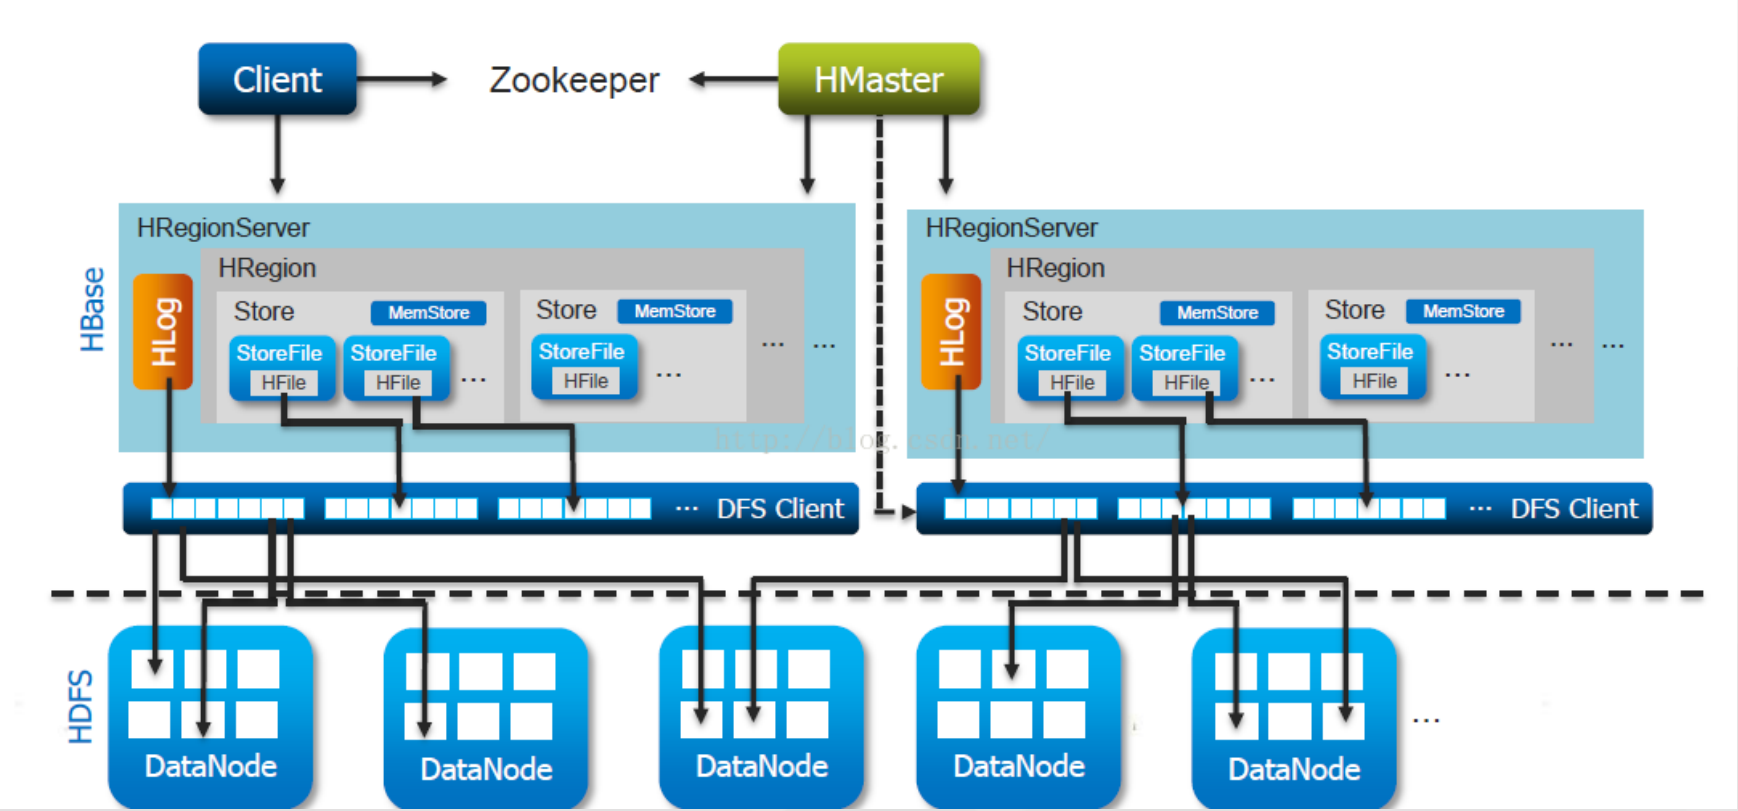
Client
Zookeeper
HMaster
HBase
HRegionServer
HRegion
HLog
Store
MemStore
StoreFile
HFile
DFS Client
HDFS
DataNode

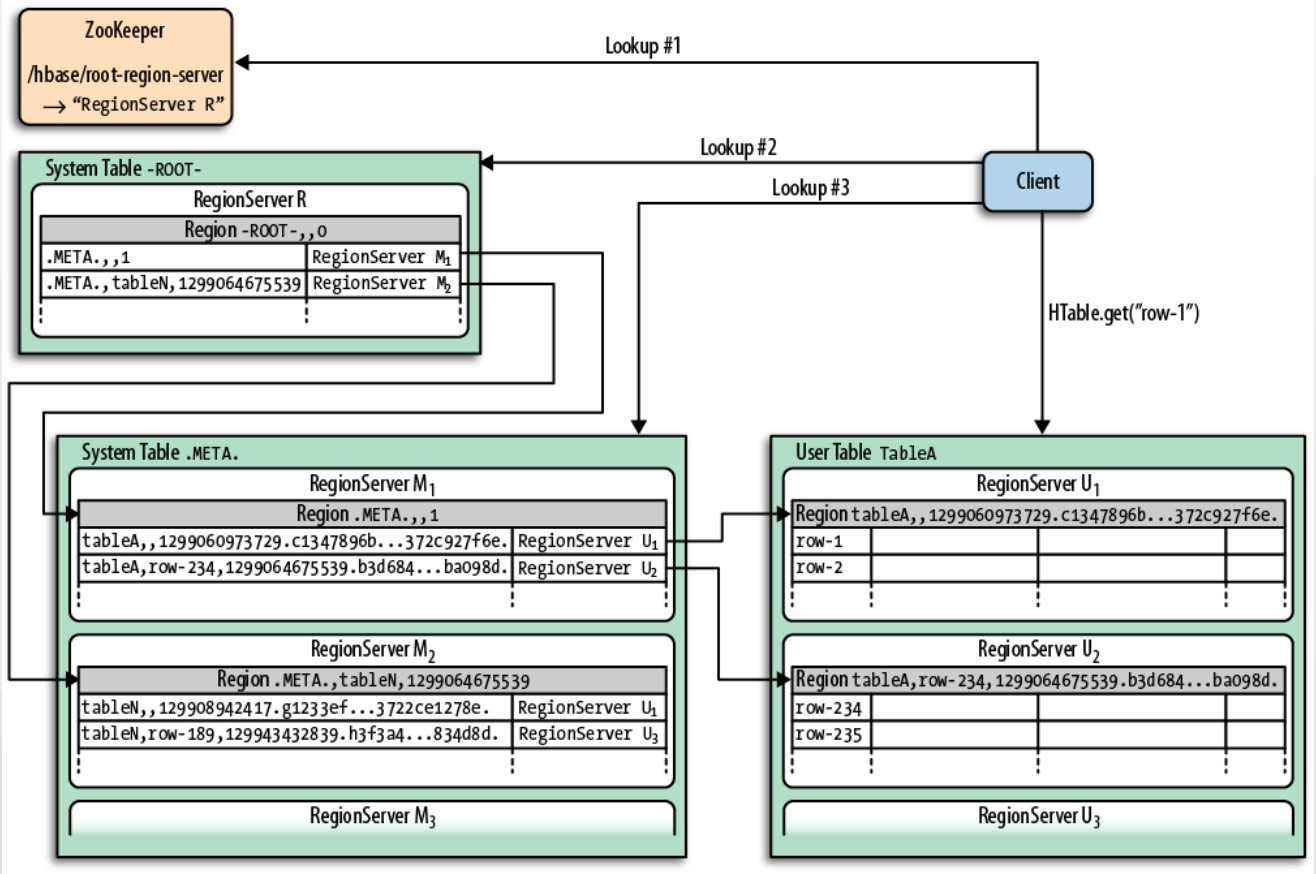
ZooKeeper
/hbase/root-region-server
→ "RegionServer R"
Lookup #1
Lookup #2
Lookup #3
Client
HTable.get("row-1")
System Table -ROOT-
RegionServer R
Region -ROOT-,,0
.META.,,1
RegionServer M1
.META.,tableN,1299064675539
RegionServer M2
System Table .META.
RegionServer M1
Region .META.,,1
tableA,,1299060973729.c1347896b...372c927f6e.
RegionServer U1
tableA,row-234,1299064675539.b3d684...ba098d.
RegionServer U2
RegionServer M2
Region .META.,tableN,1299064675539
tableN,,129908942417.g1233ef...3722ce1278e.
RegionServer U1
tableN,row-189,129943432839.h3f3a4...834d8d.
RegionServer U3
RegionServer M3
User Table TableA
RegionServer U1
Region tableA,,1299060973729.c1347896b...372c927f6e.
row-1
row-2
RegionServer U2
Region tableA,row-234,1299064675539.b3d684...ba098d.
row-234
row-235
RegionServer U3

Preliminary Topics
事前準備

JVM


Java
1.0 1996
1.2 1998
1.3 2000
1.4 2002
5.0 2004
6.0 2006
7.0 2011
8.0 2013
Scala
2001
1.0 2003
2.0 2006
2.10 2013
2003
1.0 2007
2.0 2012
FANTOM
2005
THOMPSON 2008
1.0.64 2012
Clojure
2007
1.0 2009
1.4 2012
ceylon
2010
M1 2011
M1 2012
Kotlin
2011
M1 2012
M5 2013
Xtend
1.0 2011
2.3 2012

Big Three

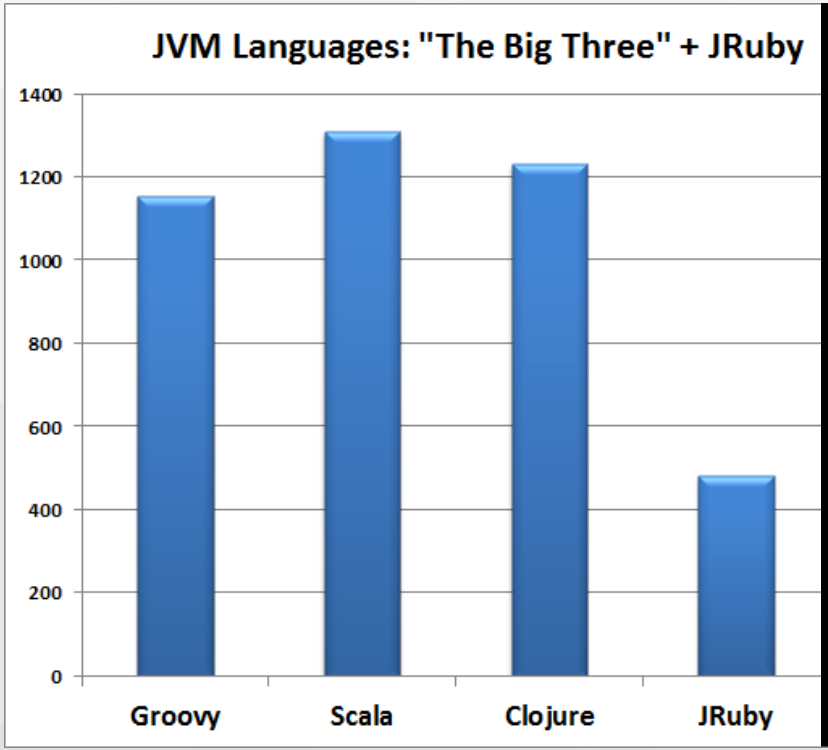
JVM Languages: "The Big Three" + JRuby
1400
1200
1000
800
600
400
200
0
Groovy
Scala
Clojure
JRuby

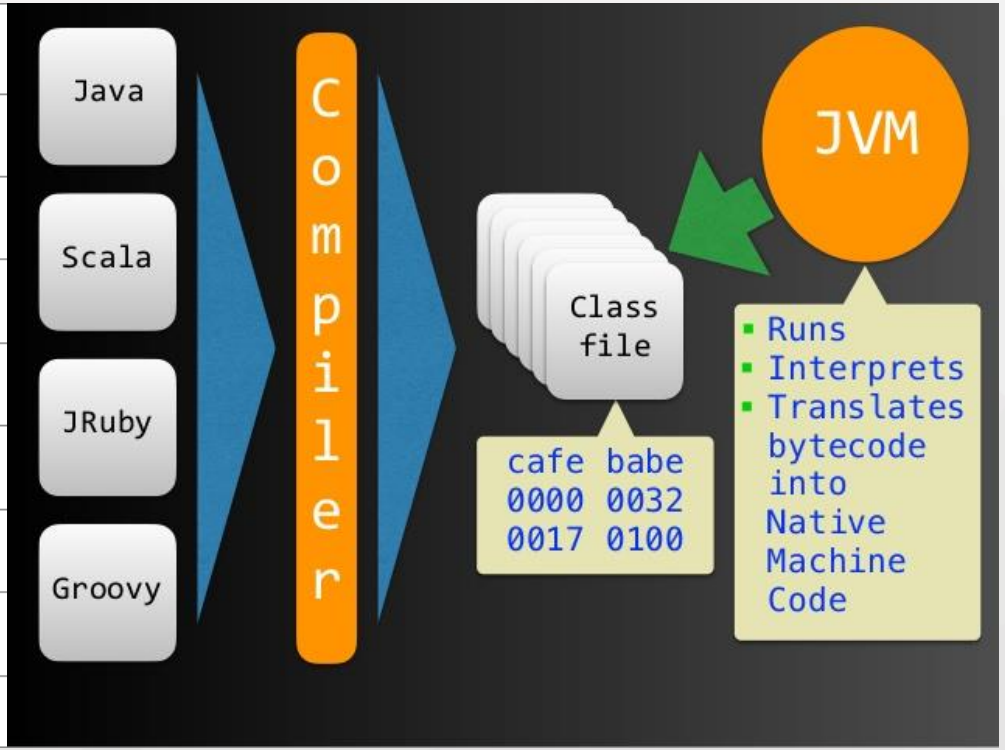
Java
Scala
JRuby
Groovy
Compiler
Class file
cafe babe
0000 0032
0017 0100
JVM
Runs
Interprets
Translates bytecode into Native Machine Code


Scala
def products = orders.flatMap(o => o.products)
public List<Product>getProducts(){
List<Product> products = new ArrayList<Product>();
for (Order order : order) {
products.addAll(order.getProducts());
}
return products;
}
Java

BUILD TOOL – SBT

http://www.scala-sbt.org/

```
name := "hello world"

version := "0.0.1"

scalaVersion := "2.11.1"

resolvers ++= Seq (
Resolver.mavenLocal,
Resolver.sonatypeRepo ("releases"),
Resolver.typesafeRepo ("releases")
)

libraryDependencies ++=
Seq ("org.scala-lang" % "scala-compiler" % "2.11.1")

addSbtPlugin ("com.typesafe.play" % "sbt-plugin" % "2.3.1")
```

Sbt


```xml
<project xmlns="http://maven.apache.org/POM/4.0.0" xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
         xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://maven.apache.org/maven-v4_0_0.xsd">
    <modelVersion>4.0.0</modelVersion>
    <groupId>info.solidsoft.rnd</groupId>
    <artifactId>spock-10-groovy-24-gradle-maven</artifactId>
    <version>0.0.1-SNAPSHOT</version>
    <properties>
        <project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
        <surefire.version>2.18.1</surefire.version>
    </properties>
    <build>
        <plugins>
            <plugin>
                <groupId>org.codehaus.gmavenplus</groupId>
                <artifactId>gmavenplus-plugin</artifactId>
                <version>1.4</version>
                <executions>
                    <execution>
                        <goals>
                            <goal>compile</goal>
                            <goal>testCompile</goal>
                        </goals>
                    </execution>
                </executions>
            </plugin>
            <plugin>
                <artifactId>maven-surefire-plugin</artifactId>
                <version>${surefire.version}</version>
                <configuration>
                    <includes>
                        <include>**/*Spec.java</include> <!-- Yes, .java extension -->
                        <include>**/*Test.java</include> <!-- Just in case having "normal" JUnit tests -->
                    </includes>
                </configuration>
            </plugin>
        </plugins>
    </build>
    <dependencies>
        <dependency>
            <groupId>org.codehaus.groovy</groupId>
            <artifactId>groovy-all</artifactId>
            <version>2.4.1</version>
        </dependency>
        <dependency>
            <groupId>org.spockframework</groupId>
            <artifactId>spock-core</artifactId>
            <version>1.0-groovy-2.4</version>
            <scope>test</scope>
        </dependency>
    </dependencies>
</project>
```

pom.xml

maven

```
apply plugin: 'groovy'

group = "info.solidsoft.rnd"
version = "0.0.1-SNAPSHOT"

repositories {
    mavenCentral()
}

dependencies {
    compile 'org.codehaus.groovy:groovy-all:2.4.1'

    testCompile 'org.spockframework:spock-core:1.0-groovy-2.4'
}
```

build.gradle

```
rootProject.name = 'spock-10-groovy-24-gradle-maven'
```

settings.xml

gradle

Yarn

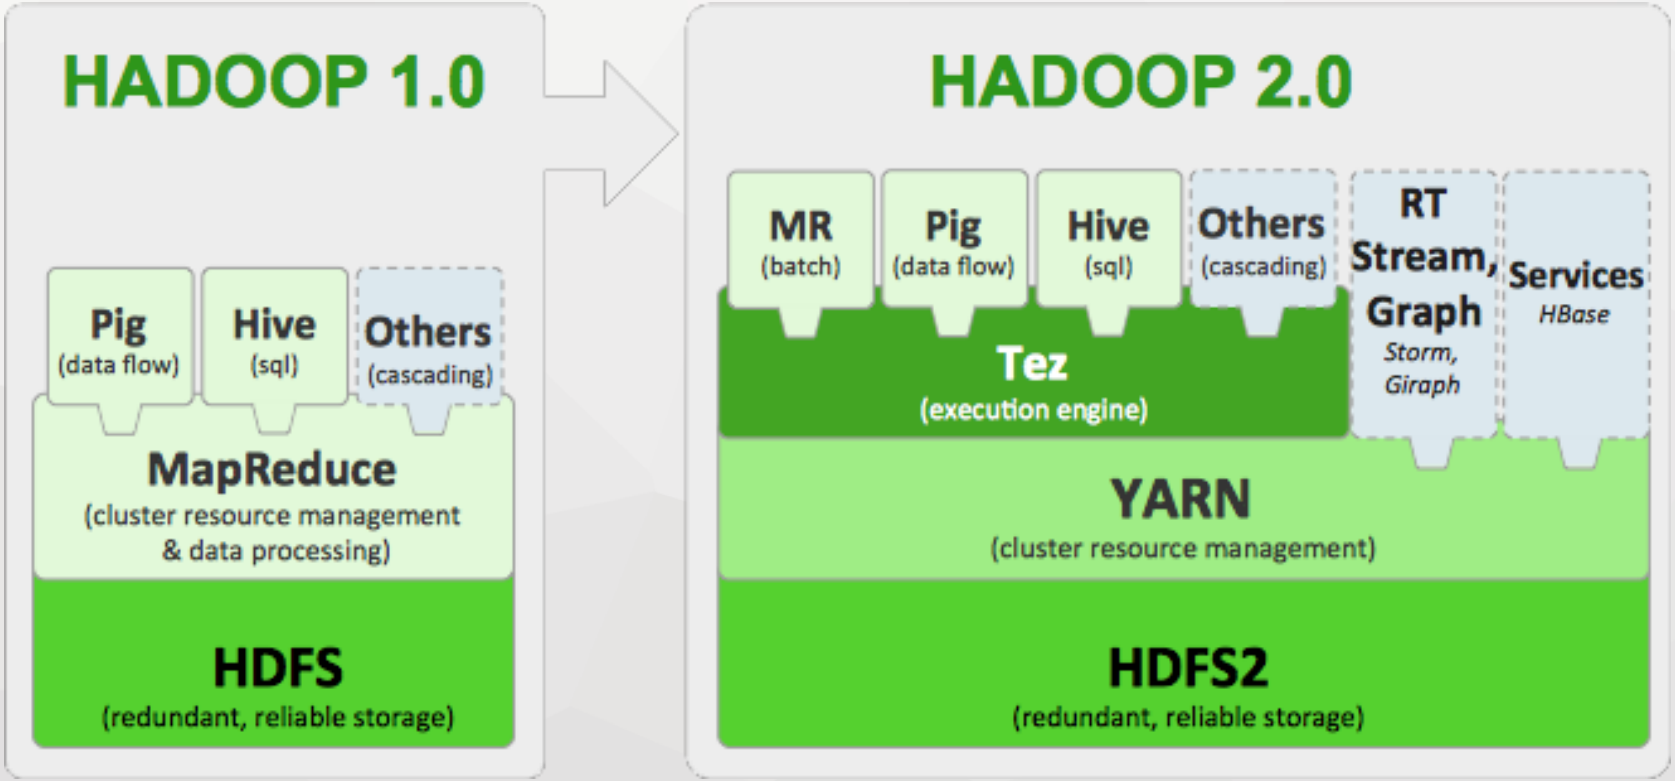
HADOOP 1.0
Pig
(data flow)
Hive
(sql)
Others
(cascading)
MapReduce
(cluster resource management
& data processing)
HDFS
(redundant, reliable storage)
HADOOP 2.0
MR
(batch)
Pig
(data flow)
Hive
(sql)
Others
(cascading)
RT
Stream,
Graph
Storm,
Giraph
Services
HBase
Tez
(execution engine)
YARN
(cluster resource management)
HDFS2
(redundant, reliable storage)

Yarn

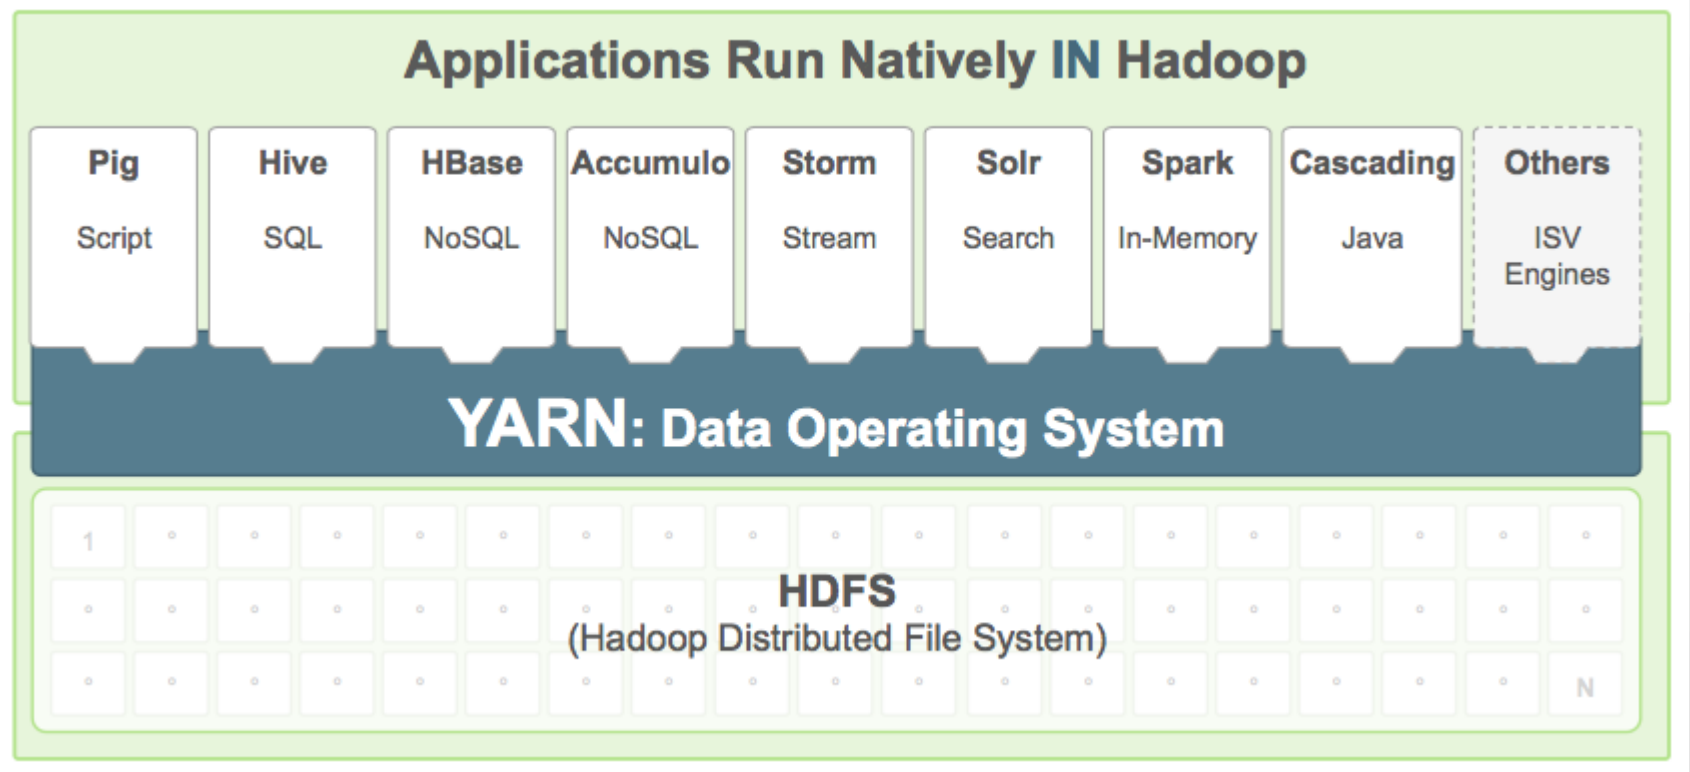
Applications Run Natively IN Hadoop
Pig
Script
Hive
SQL
HBase
NoSQL
Accumulo
NoSQL
Storm
Stream
Solr
Search
Spark
In-Memory
Cascading
Java
Others
ISV Engines
YARN: Data Operating System
1
N
HDFS
(Hadoop Distributed File System)

Yarn

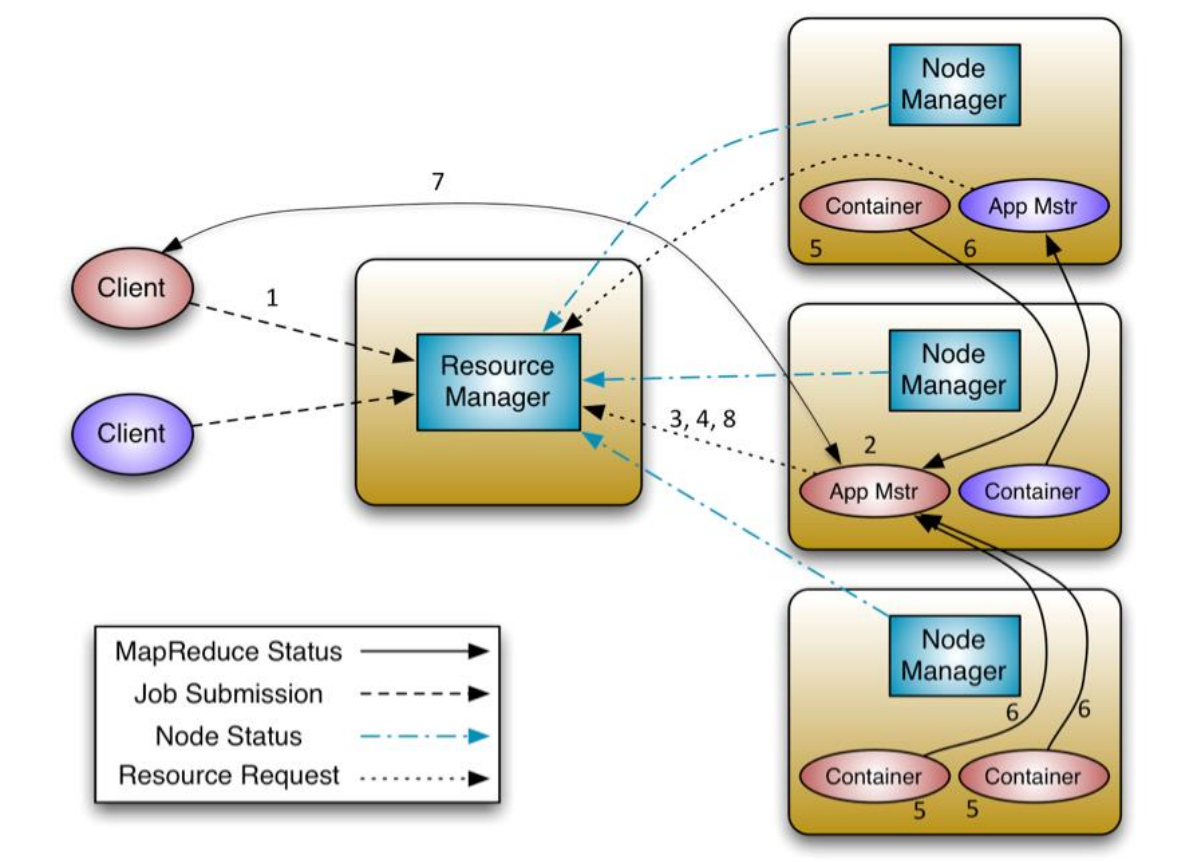

https://hadoop.apache.org/docs/stable/hadoop-yarn/hadoop-yarn-site/WritingYarnApplications.html

Mesos

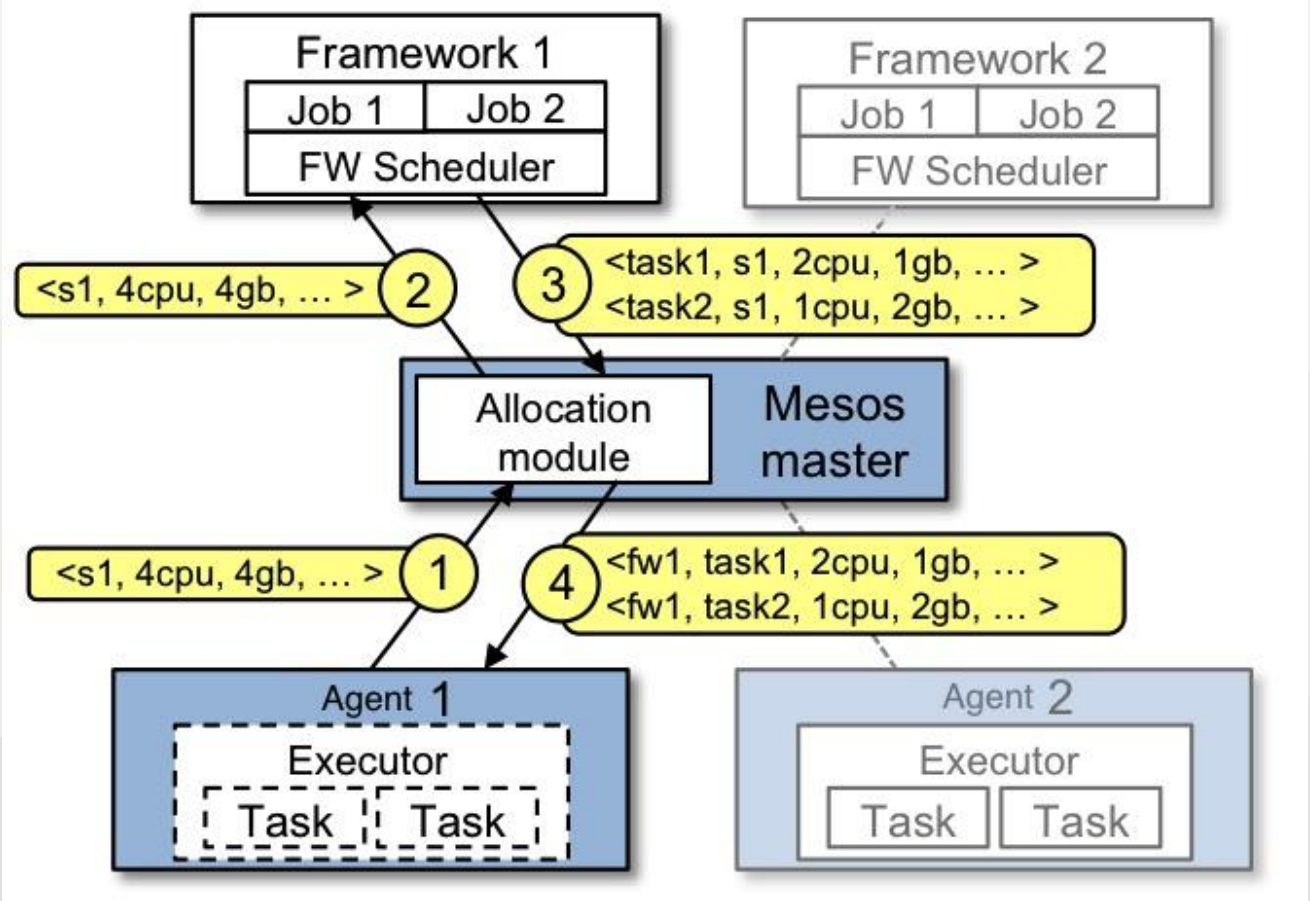
Framework 1
Job 1
Job 2
FW Scheduler
Framework 2
Job 1
Job 2
FW Scheduler
<s1, 4cpu, 4gb, … >
2
3
<task1, s1, 2cpu, 1gb, … >
<task2, s1, 1cpu, 2gb, … >
Allocation module
Mesos master
<s1, 4cpu, 4gb, … >
1
4
<fw1, task1, 2cpu, 1gb, … >
<fw1, task2, 1cpu, 2gb, … >
Agent 1
Executor
Task
Task
Agent 2
Executor
Task
Task

Spark Environment

Spark Standalone

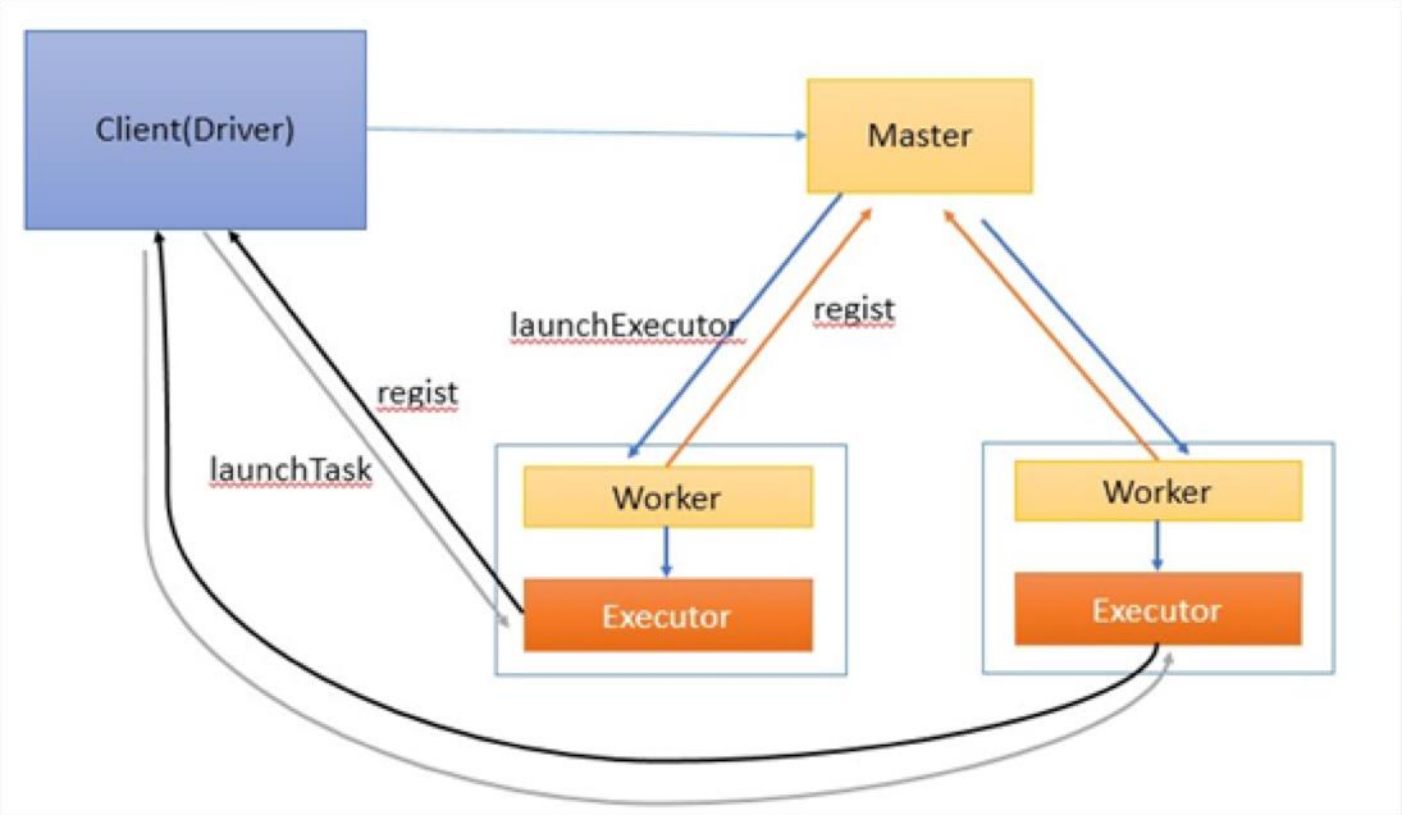

Spark Standalone HA

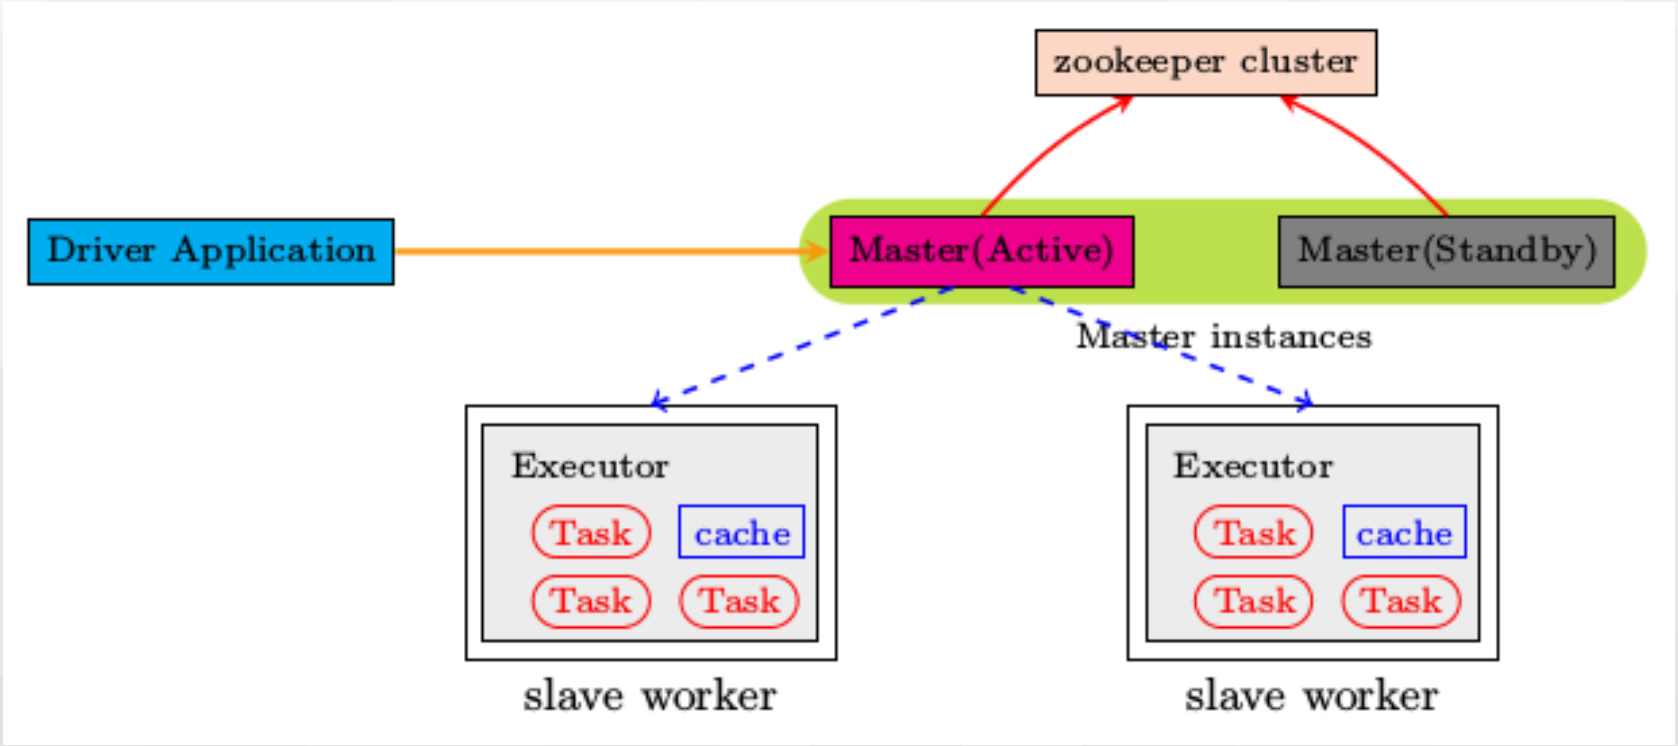

Spark Source Compile

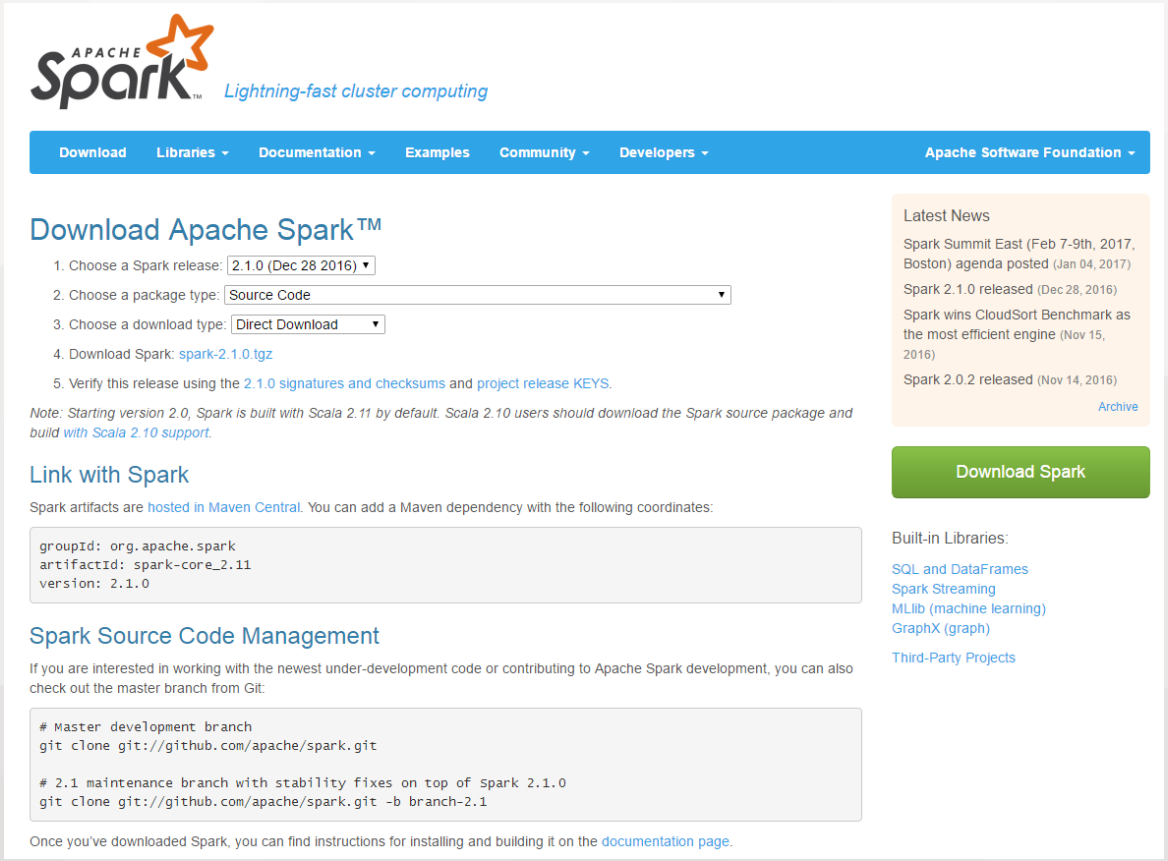

Lightning-fast cluster computing

Download | Libraries | Documentation | Examples | Community | Developers | Apache Software Foundation

Download Apache Spark™

1. Choose a Spark release: 2.1.0 (Dec 28 2016)
2. Choose a package type: Source Code
3. Choose a download type: Direct Download
4. Download Spark: spark-2.1.0.tgz
5. Verify this release using the 2.1.0 signatures and checksums and project release KEYS.

Note: Starting version 2.0, Spark is built with Scala 2.11 by default. Scala 2.10 users should download the Spark source package and build with Scala 2.10 support.

Link with Spark

Spark artifacts are hosted in Maven Central. You can add a Maven dependency with the following coordinates:

```
groupId: org.apache.spark
artifactId: spark-core_2.11
version: 2.1.0
```

Spark Source Code Management

If you are interested in working with the newest under-development code or contributing to Apache Spark development, you can also check out the master branch from Git:

```
# Master development branch
git clone git://github.com/apache/spark.git

# 2.1 maintenance branch with stability fixes on top of Spark 2.1.0
git clone git://github.com/apache/spark.git -b branch-2.1
```

Once you've downloaded Spark, you can find instructions for installing and building it on the documentation page.

Latest News

Spark Summit East (Feb 7-9th, 2017, Boston) agenda posted (Jan 04, 2017)

Spark 2.1.0 released (Dec 28, 2016)

Spark wins CloudSort Benchmark as the most efficient engine (Nov 15, 2016)

Spark 2.0.2 released (Nov 14, 2016)

Archive

Download Spark

Built-in Libraries:

SQL and DataFrames
Spark Streaming
MLlib (machine learning)
GraphX (graph)

Third-Party Projects

Spark On Yarn

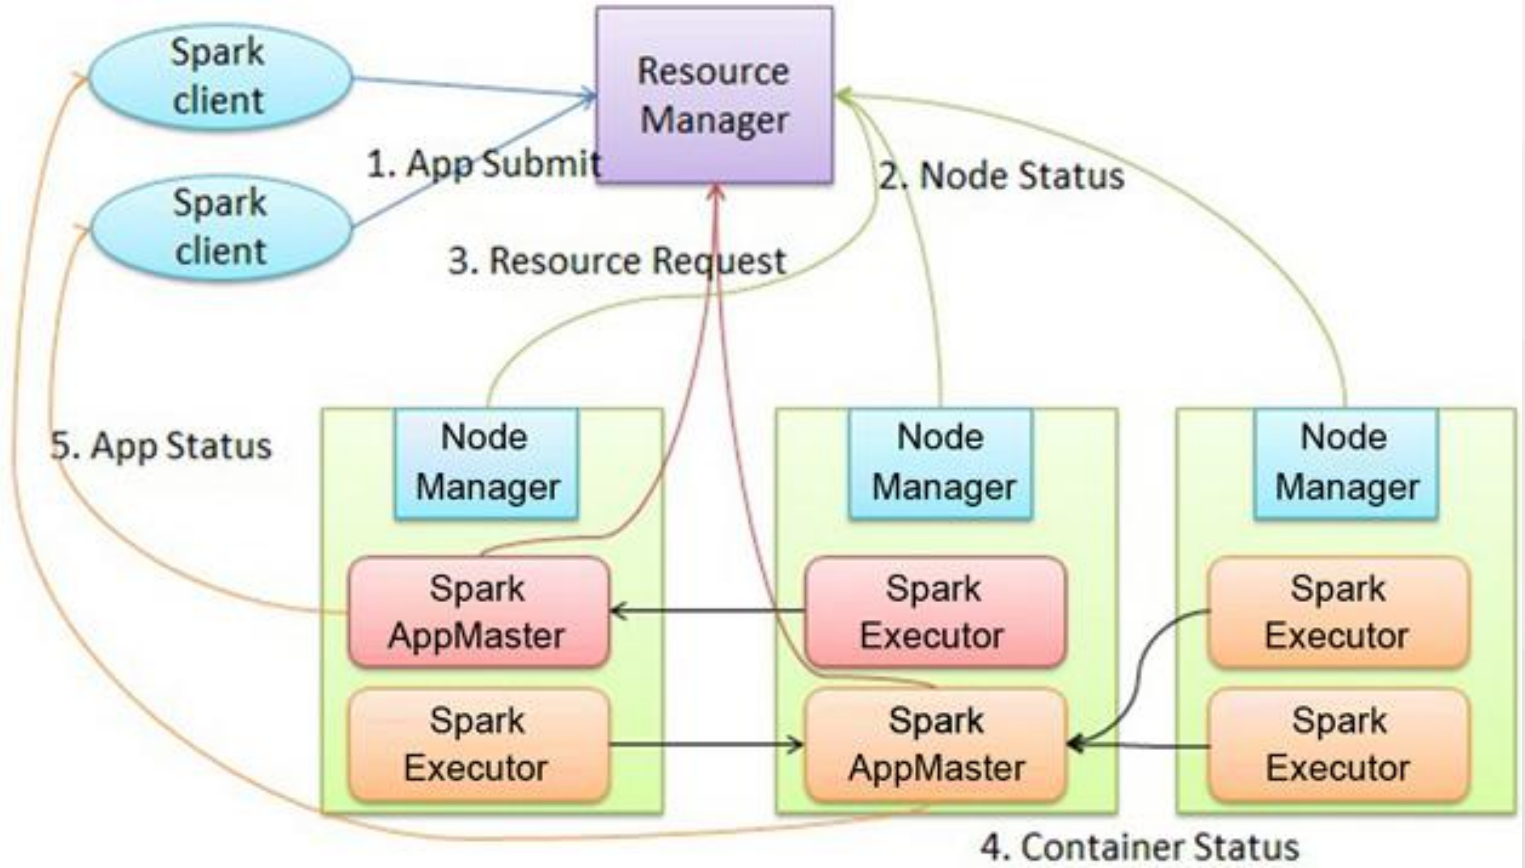

Spark On Mesos

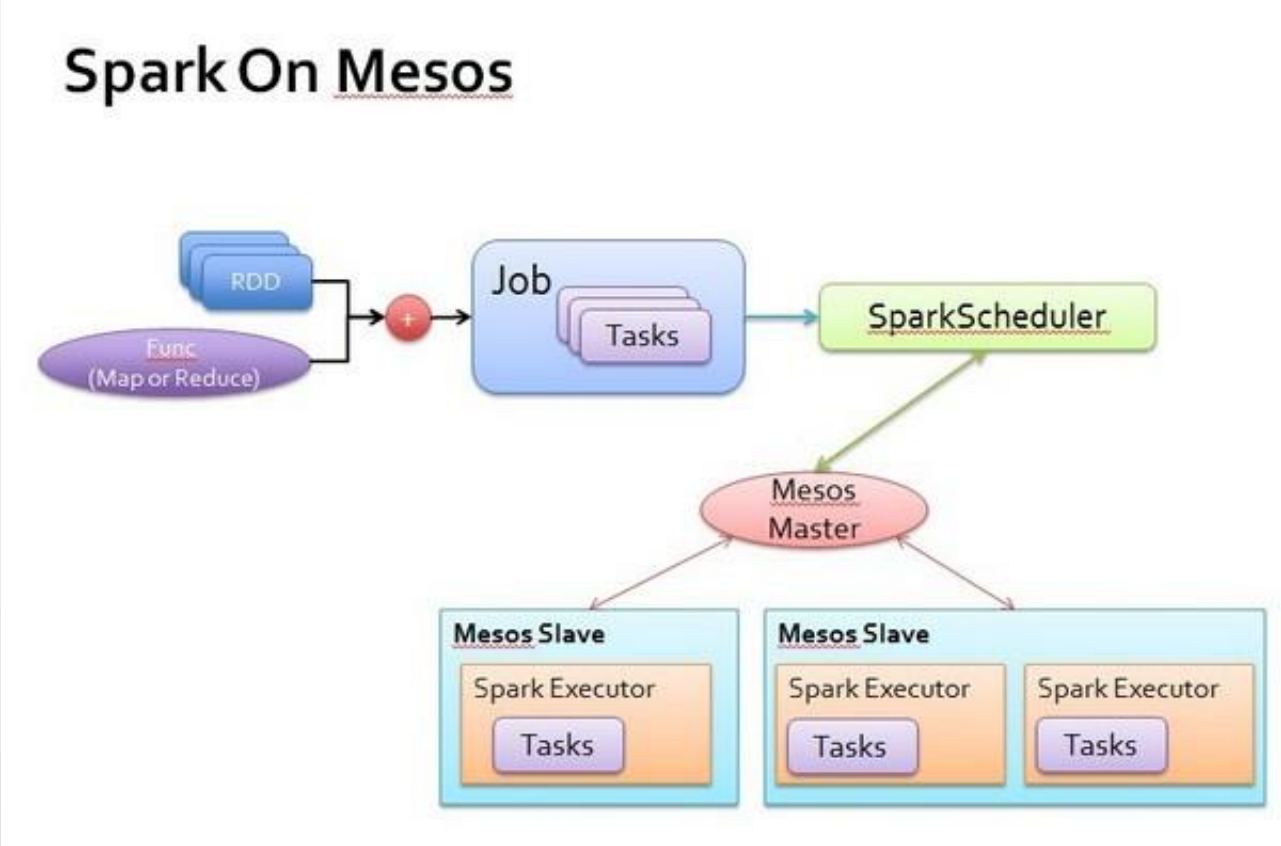

Spark Architecture

Ecosystems

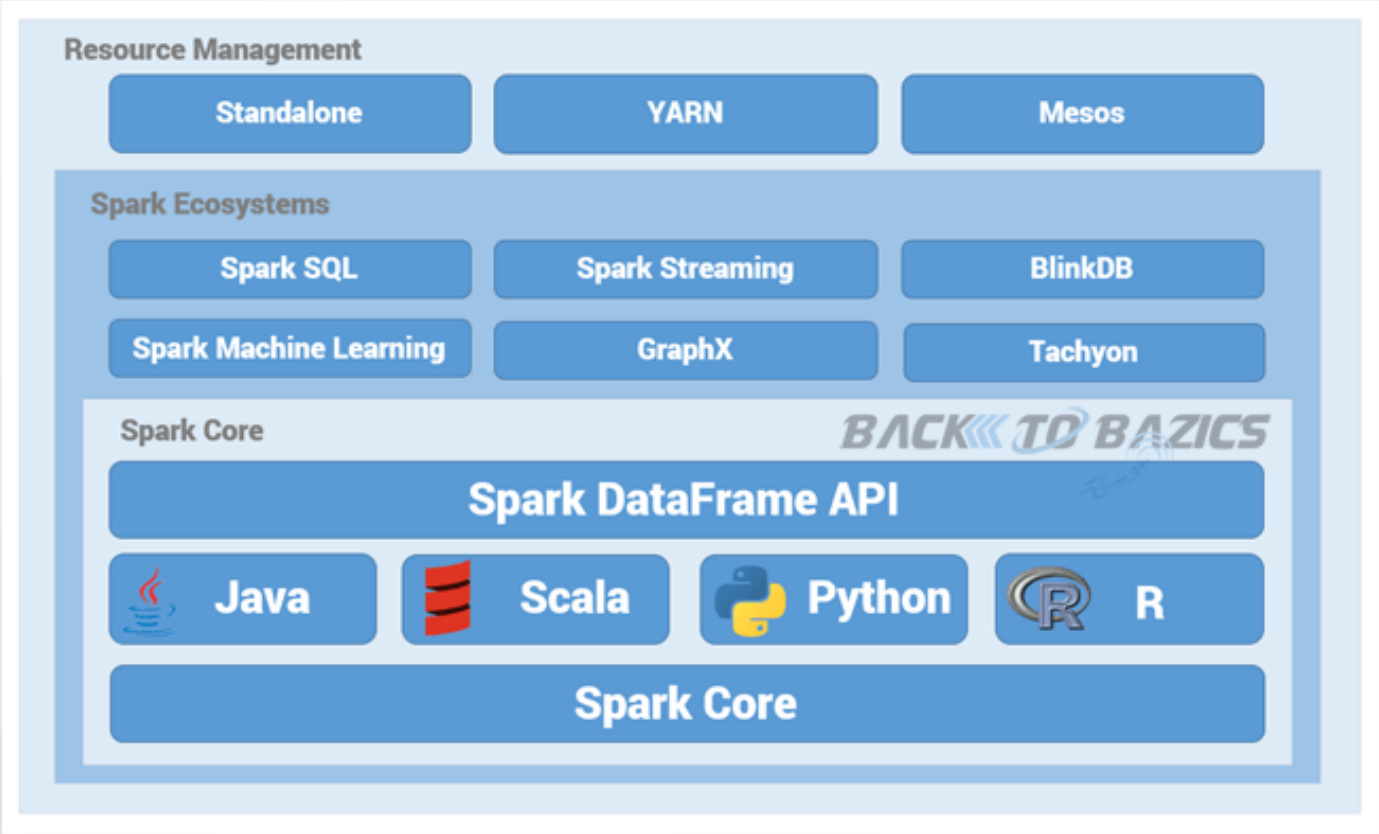
Resource Management
Standalone
YARN
Mesos
Spark Ecosystems
Spark SQL
Spark Streaming
BlinkDB
Spark Machine Learning
GraphX
Tachyon
Spark Core
BACK TO BAZICS
Spark DataFrame API
Java
Scala
Python
R
Spark Core

Architecture

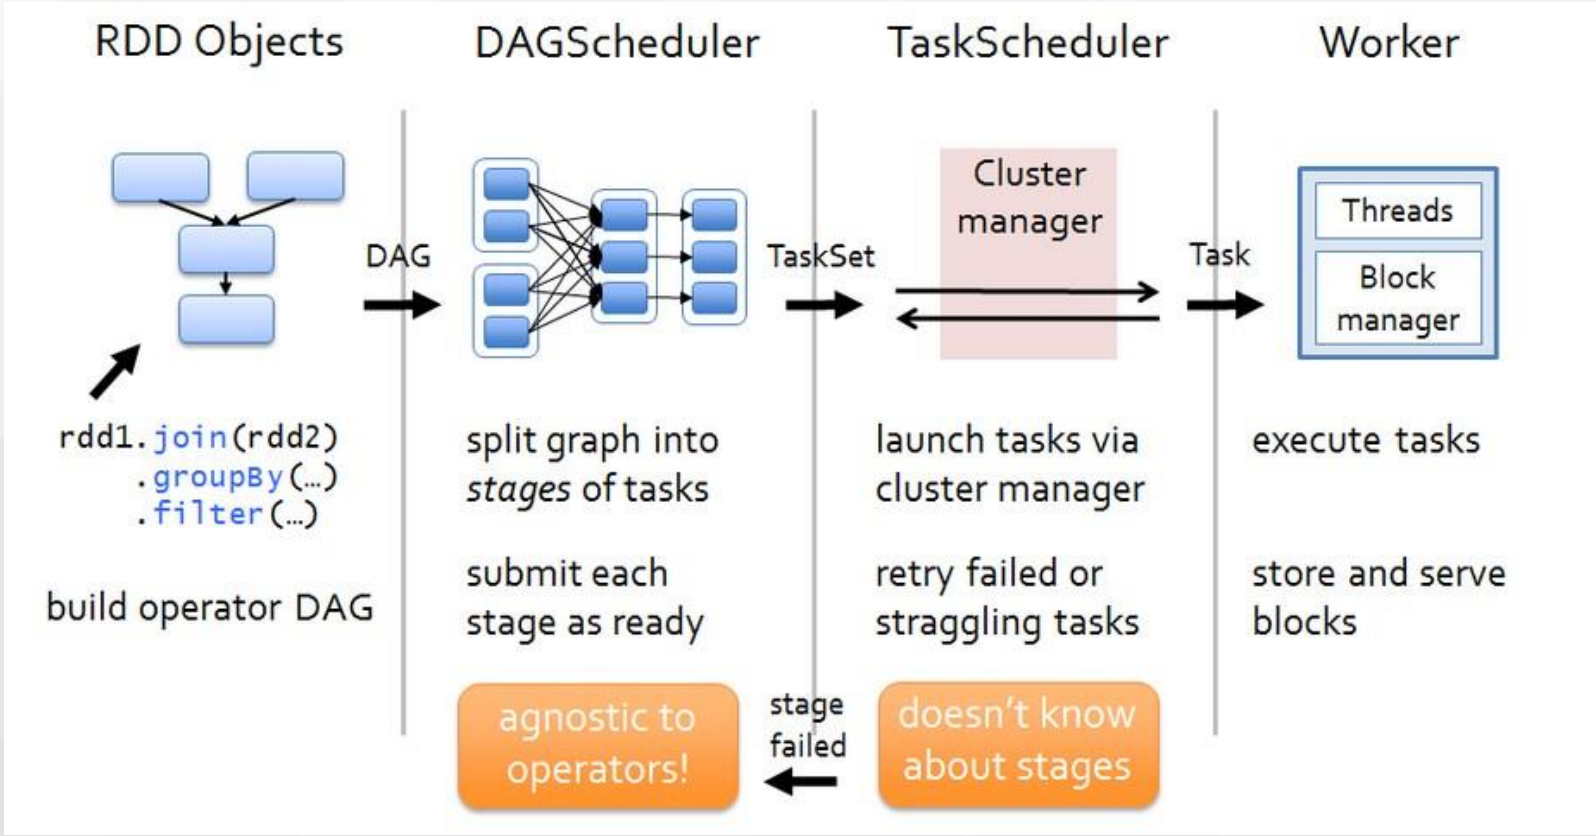

What is an RDD?

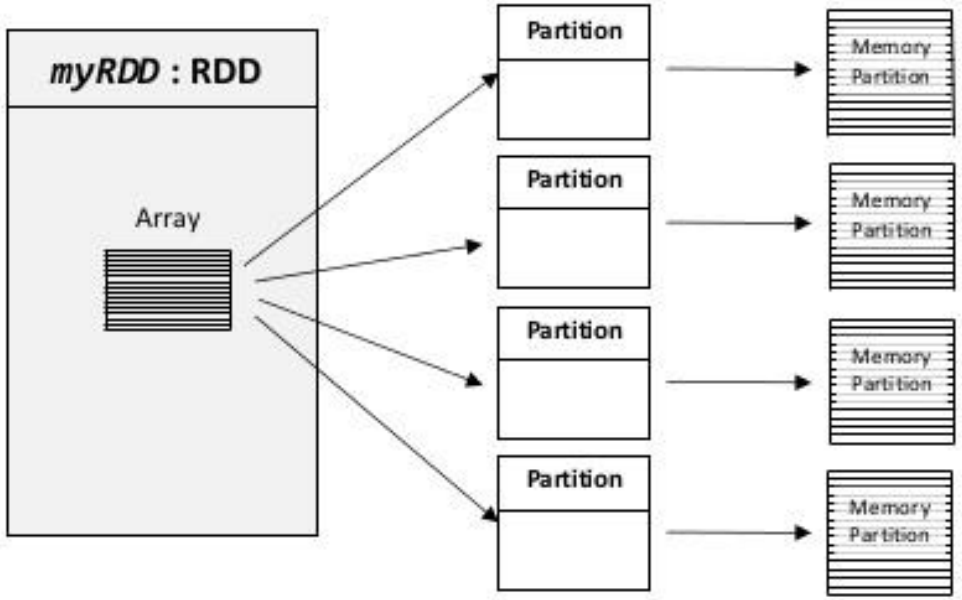

Some RDD Characteristics

- Hold references to Partition objects
- Each Partition object references a subset of your data
- Partitions are assigned to nodes on your cluster
- Each partition/split will be in RAM (by default)

Dependency Type

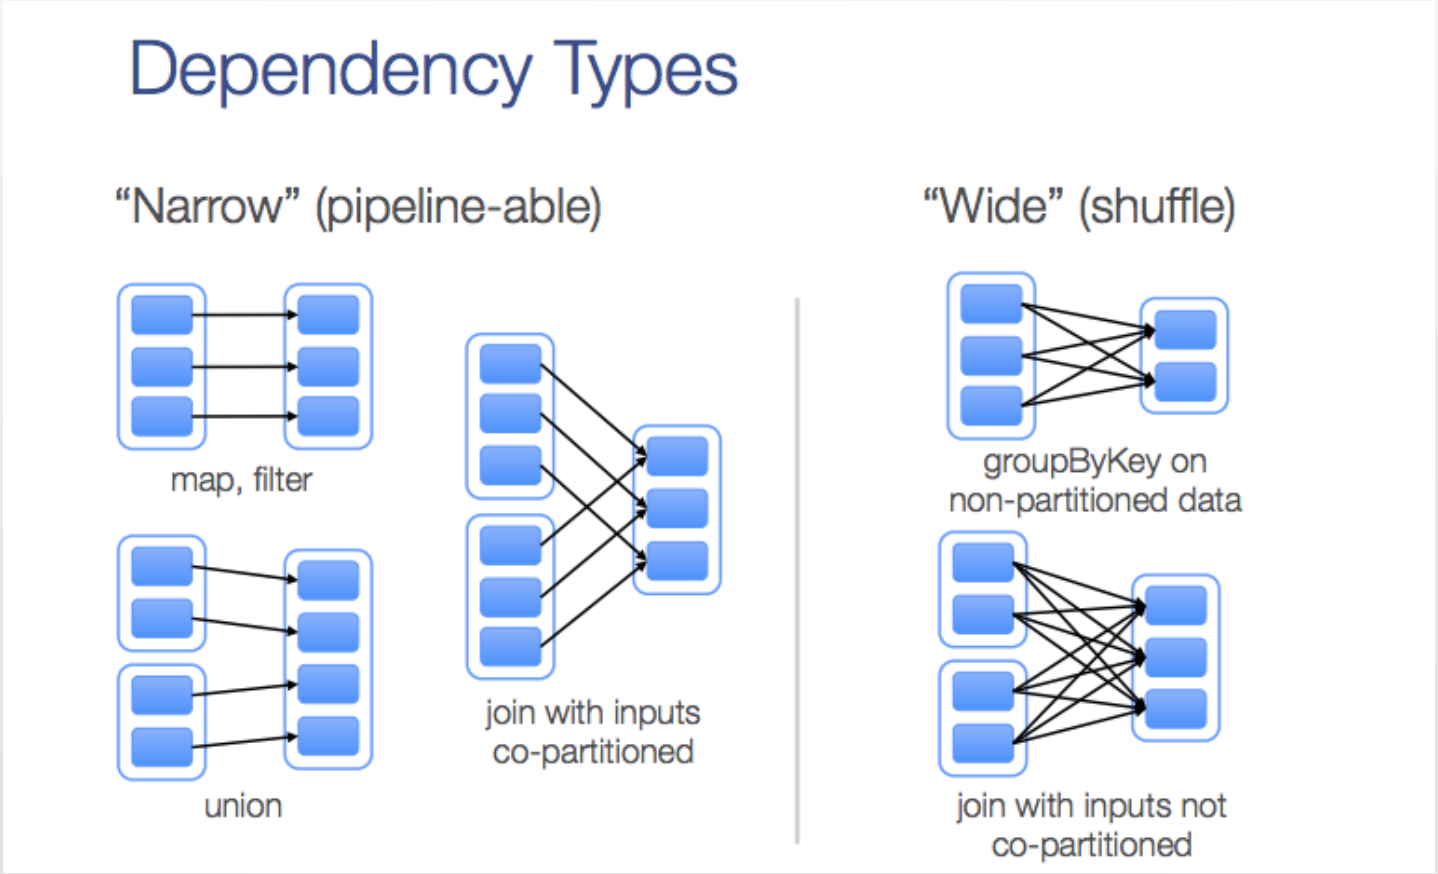

Dependency Type

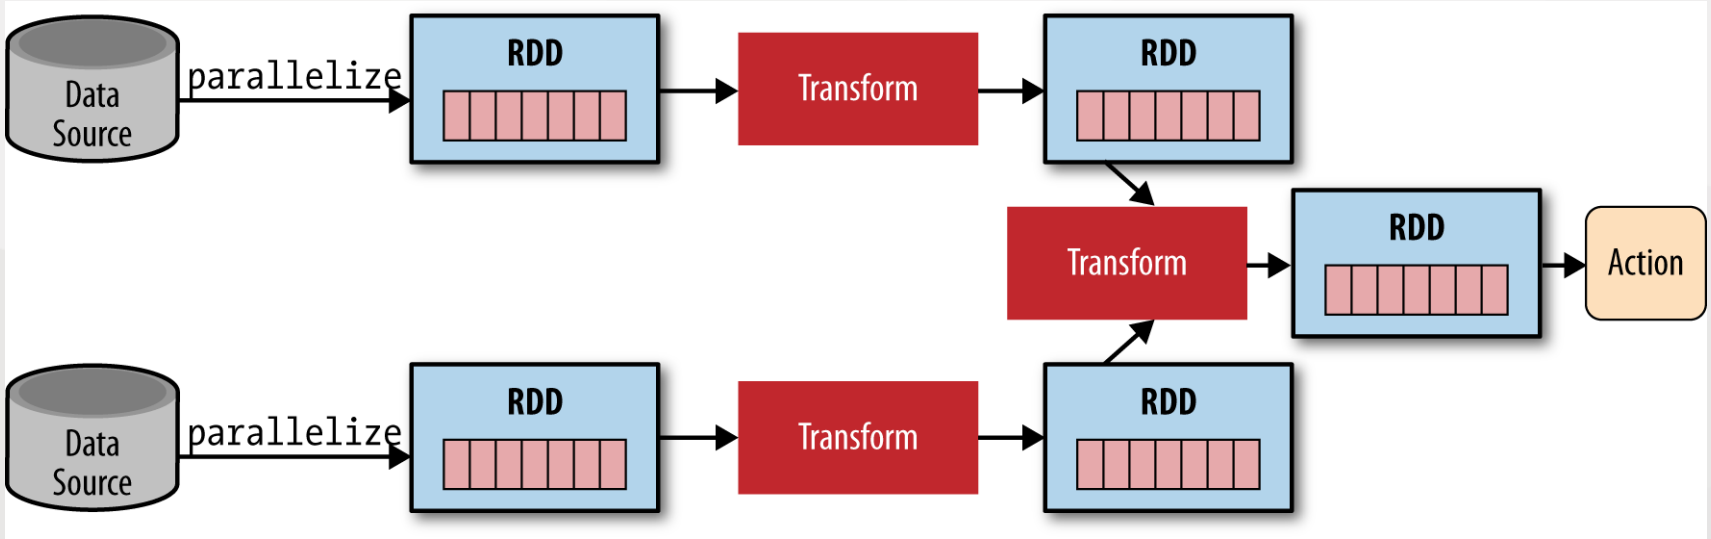

Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
cachedMsgs = messages.cache()

cachedMsgs.filter(_.contains("foo")).count
cachedMsgs.filter(_.contains("bar")).count
. . .
```

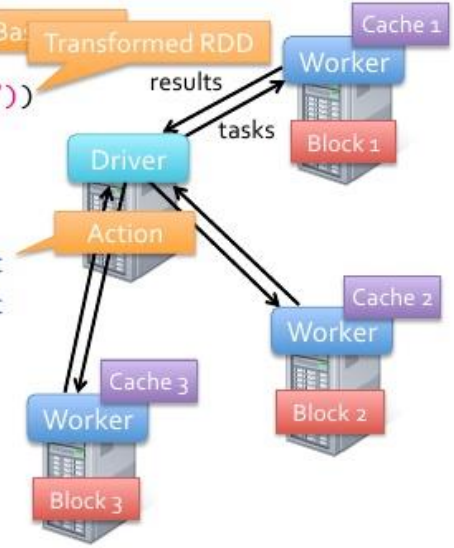

Result: scaled to 1 TB data in 5-7 sec
(vs 170 sec for on-disk data)

Driver
Executor
Application

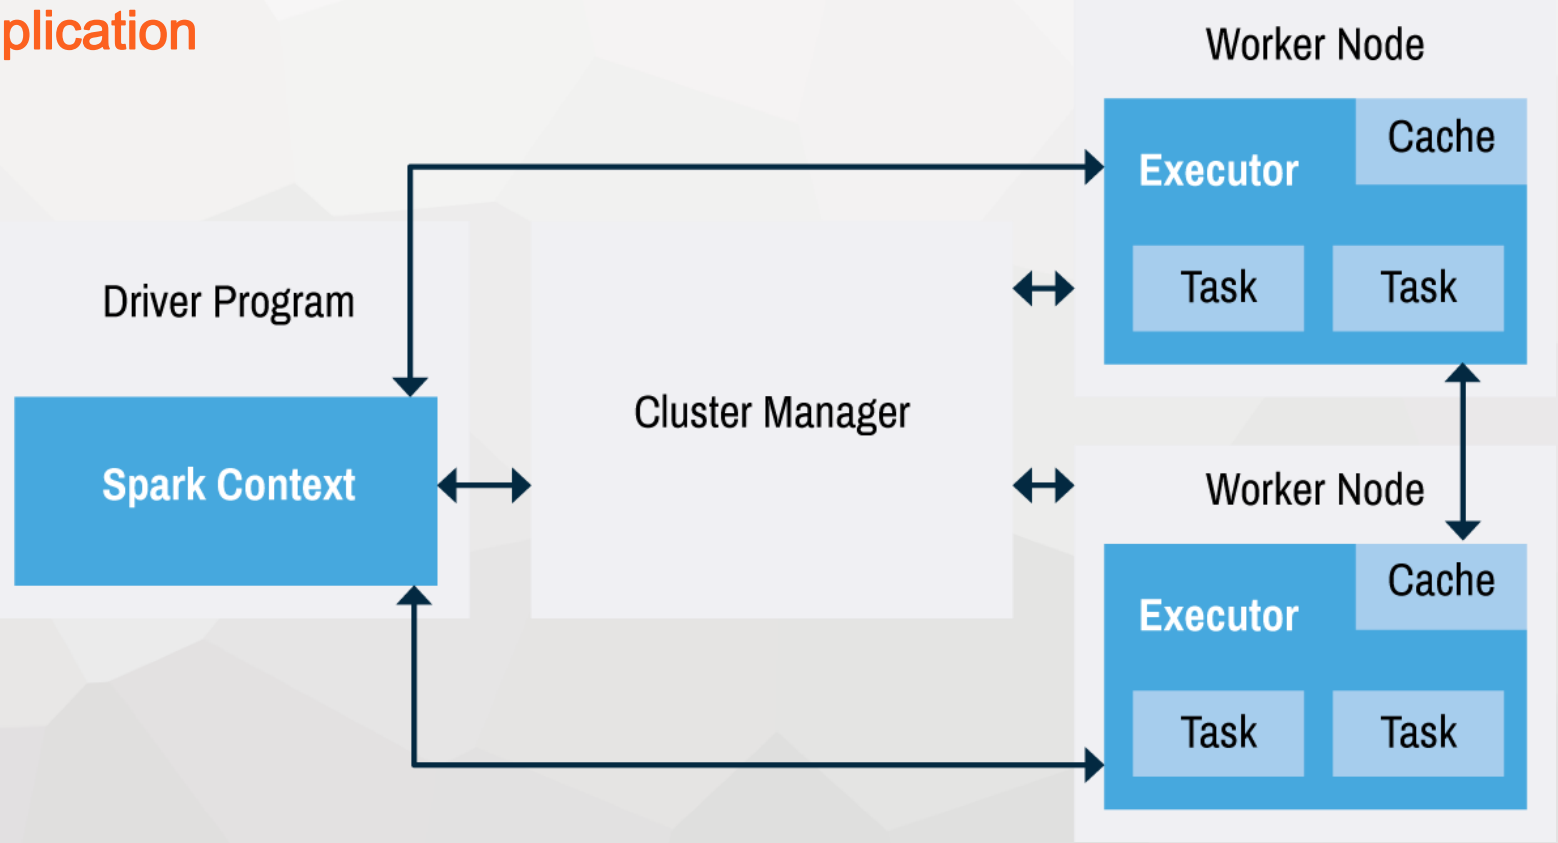

Job

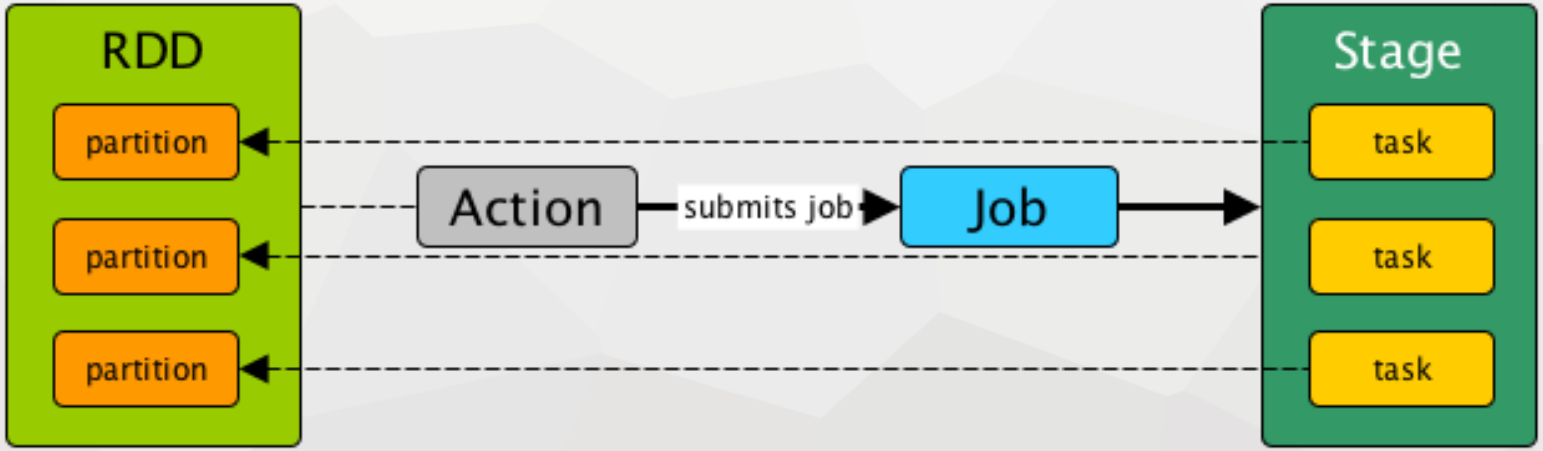
RDD
partition
partition
partition
Action
submits job
Job
Stage
task
task
task

Stage

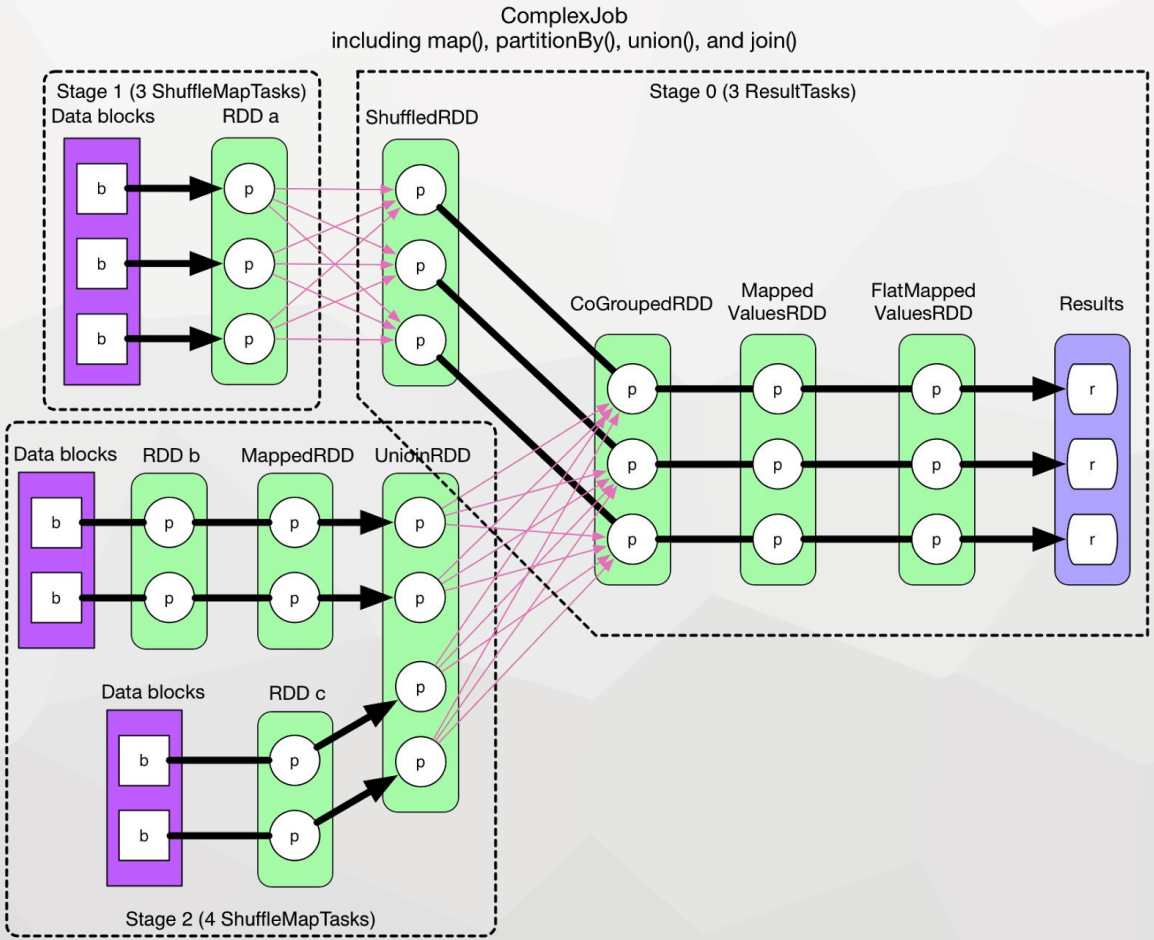
ComplexJob
including map(), partitionBy(), union(), and join()
Stage 1 (3 ShuffleMapTasks)
Data blocks
RDD a
Stage 0 (3 ResultTasks)
ShuffledRDD
CoGroupedRDD
Mapped ValuesRDD
FlatMapped ValuesRDD
Results
Data blocks
RDD b
MappedRDD
UnionRDD
Data blocks
RDD c
Stage 2 (4 ShuffleMapTasks)

Persist & Cache

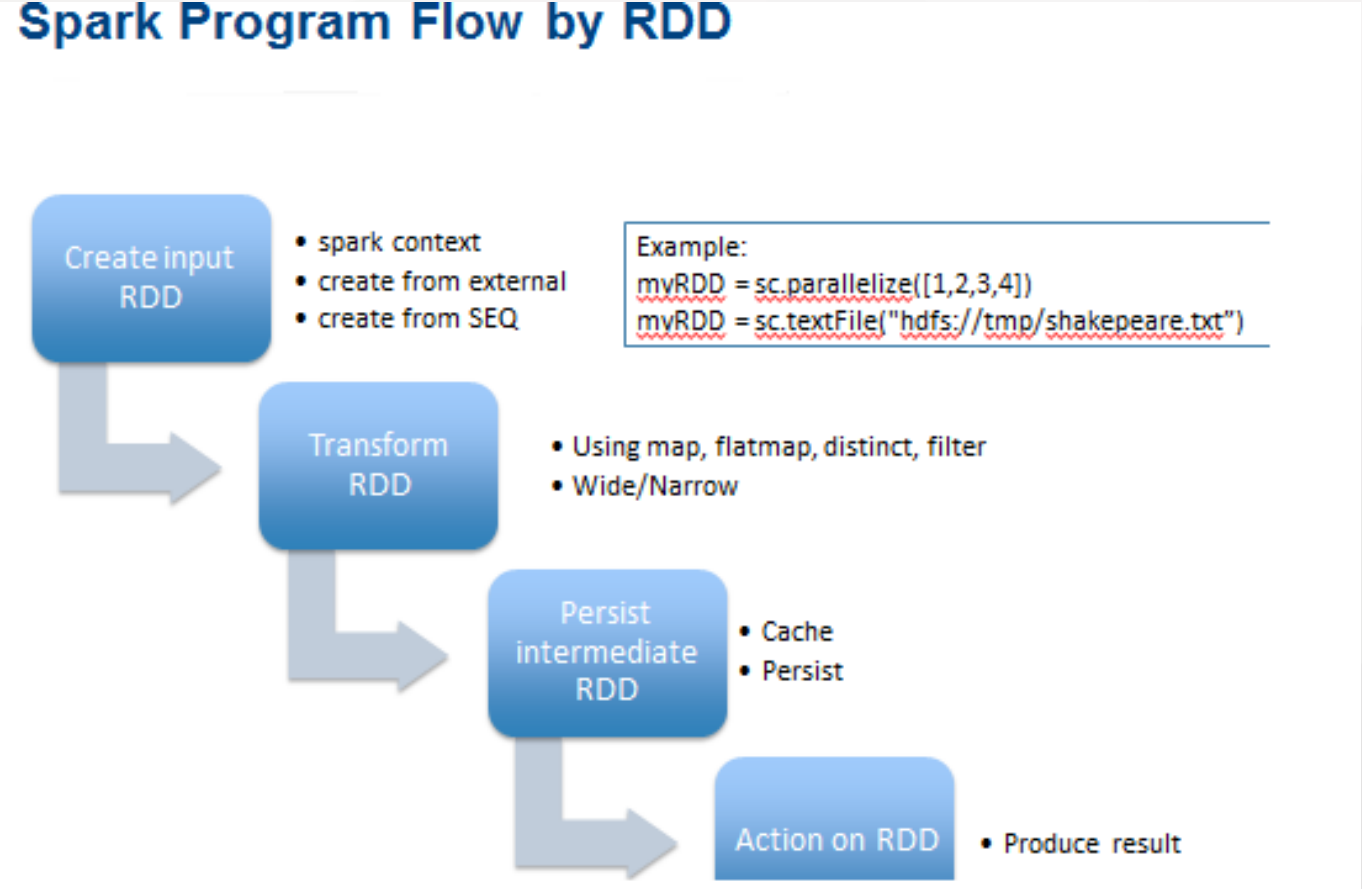

Spark SQL

Hive

- Data warehousing package built on top of Hadoop
- Bringing structure to unstructured data
- Query petabytes of data with HiveQL
- Schema on read

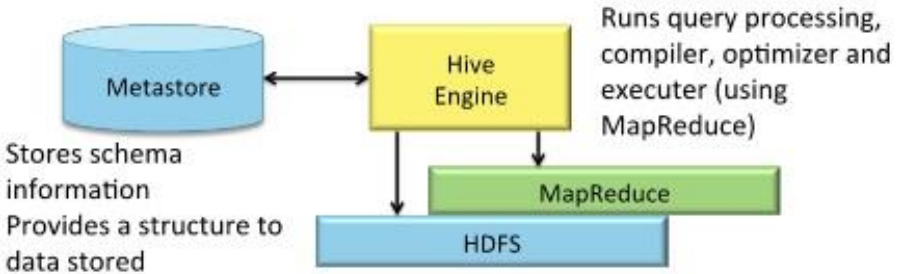

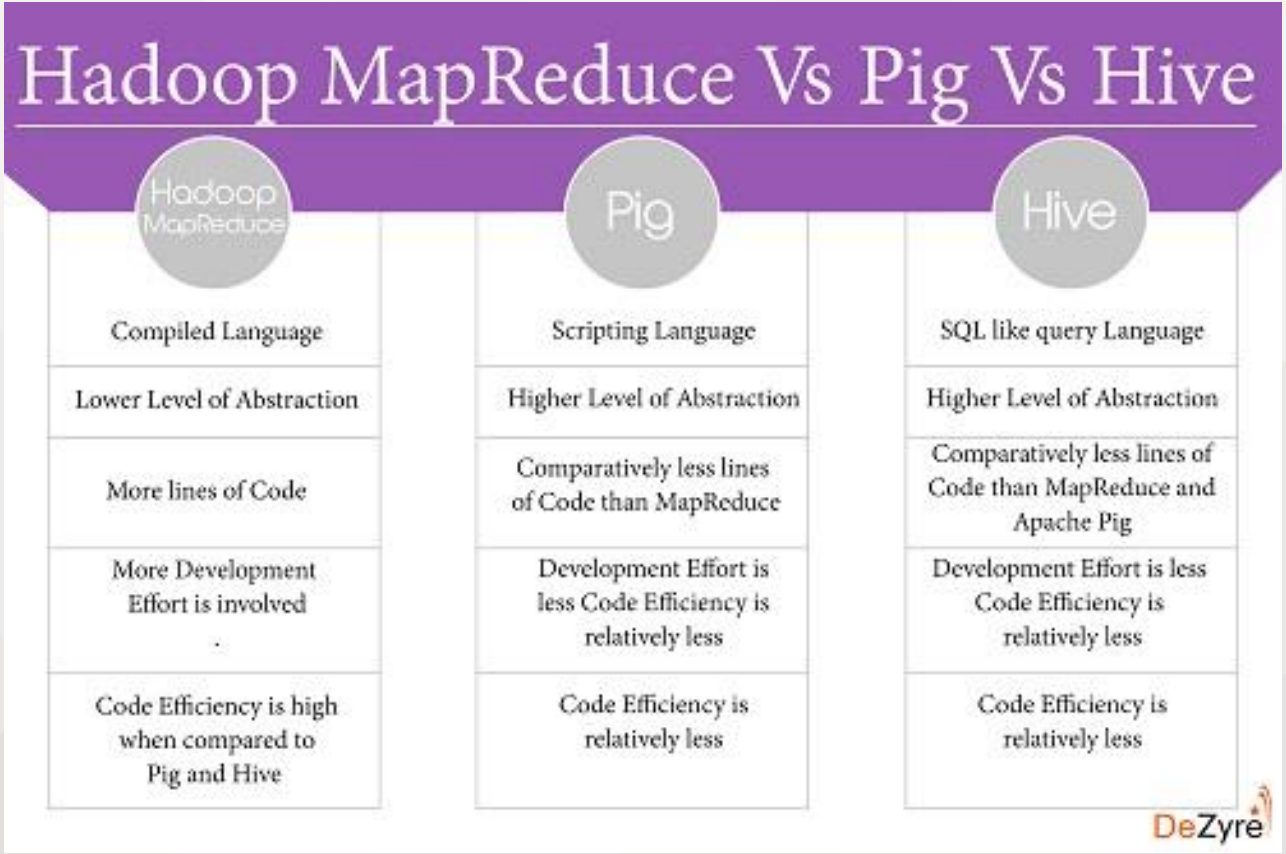
Hadoop MapReduce Vs Pig Vs Hive
Hadoop MapReduce
Pig
Hive
Compiled Language
Lower Level of Abstraction
More lines of Code
More Development Effort is involved
.
Code Efficiency is high when compared to Pig and Hive
Scripting Language
Higher Level of Abstraction
Comparatively less lines of Code than MapReduce
Development Effort is less Code Efficiency is relatively less
Code Efficiency is relatively less
SQL like query Language
Higher Level of Abstraction
Comparatively less lines of Code than MapReduce and Apache Pig
Development Effort is less Code Efficiency is relatively less
Code Efficiency is relatively less
DeZyre

The Right SQL Engine for the Use Case

BI and SQL Analytics

Batch Processing

Spark Developers

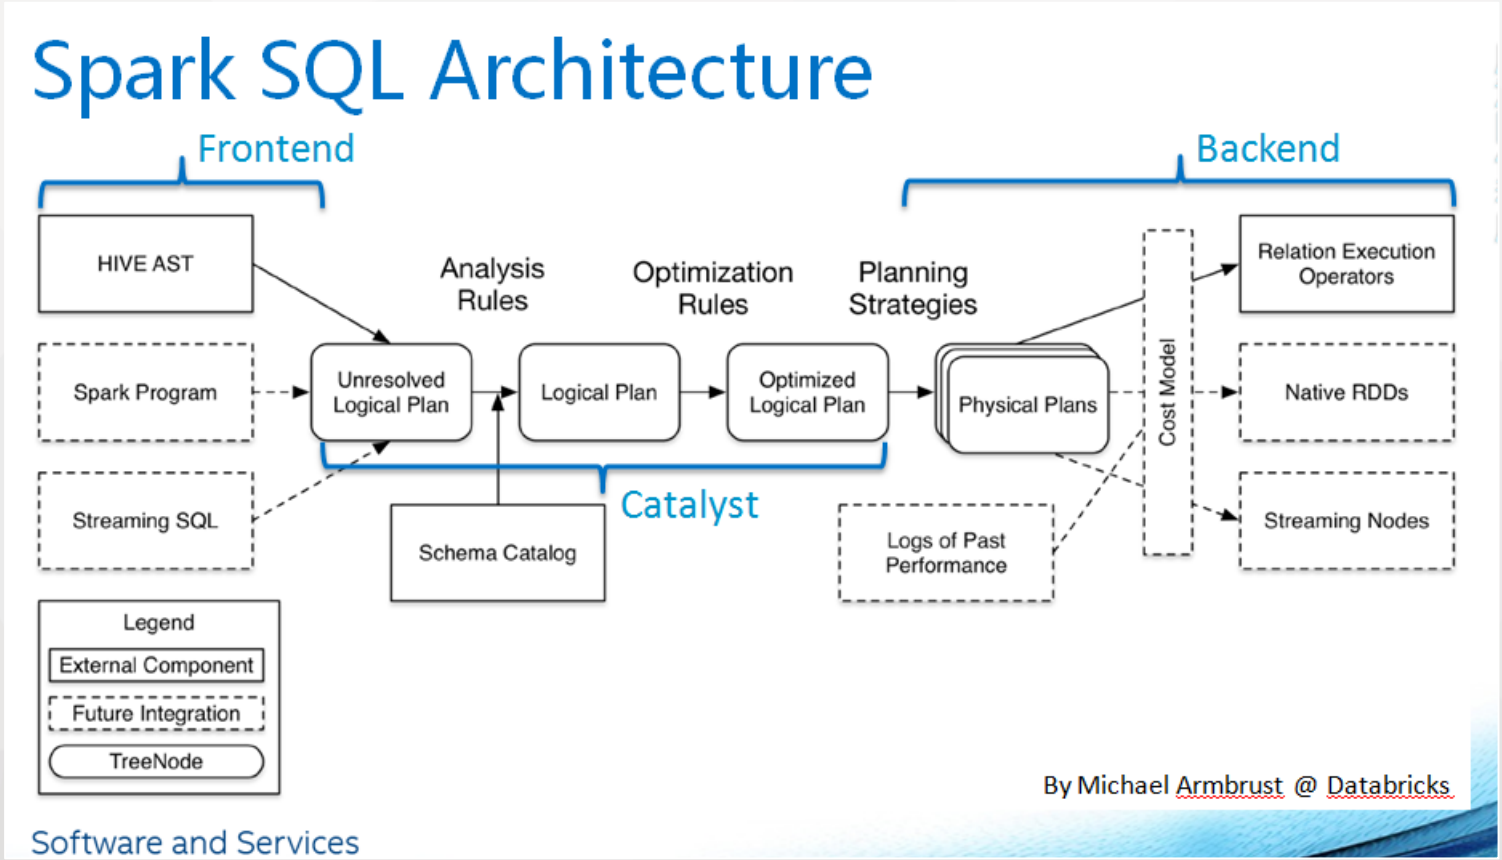
Spark SQL Architecture
Frontend
Backend
HIVE AST
Spark Program
Streaming SQL
Analysis Rules
Optimization Rules
Planning Strategies
Unresolved Logical Plan
Logical Plan
Optimized Logical Plan
Physical Plans
Cost Model
Relation Execution Operators
Native RDDs
Streaming Nodes
Catalyst
Schema Catalog
Logs of Past Performance
Legend
External Component
Future Integration
TreeNode
By Michael Armbrust @ Databricks
Software and Services

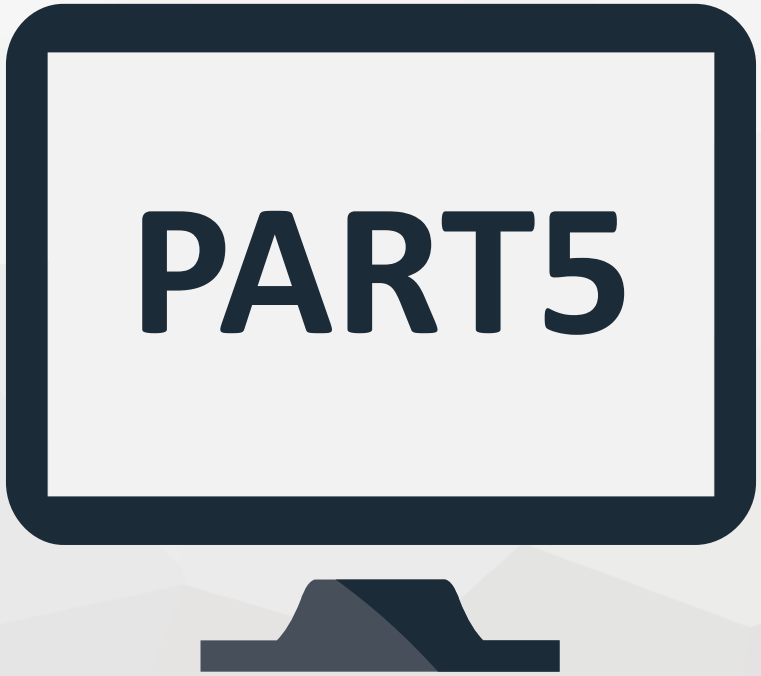

Spark Streaming

Ecosystems

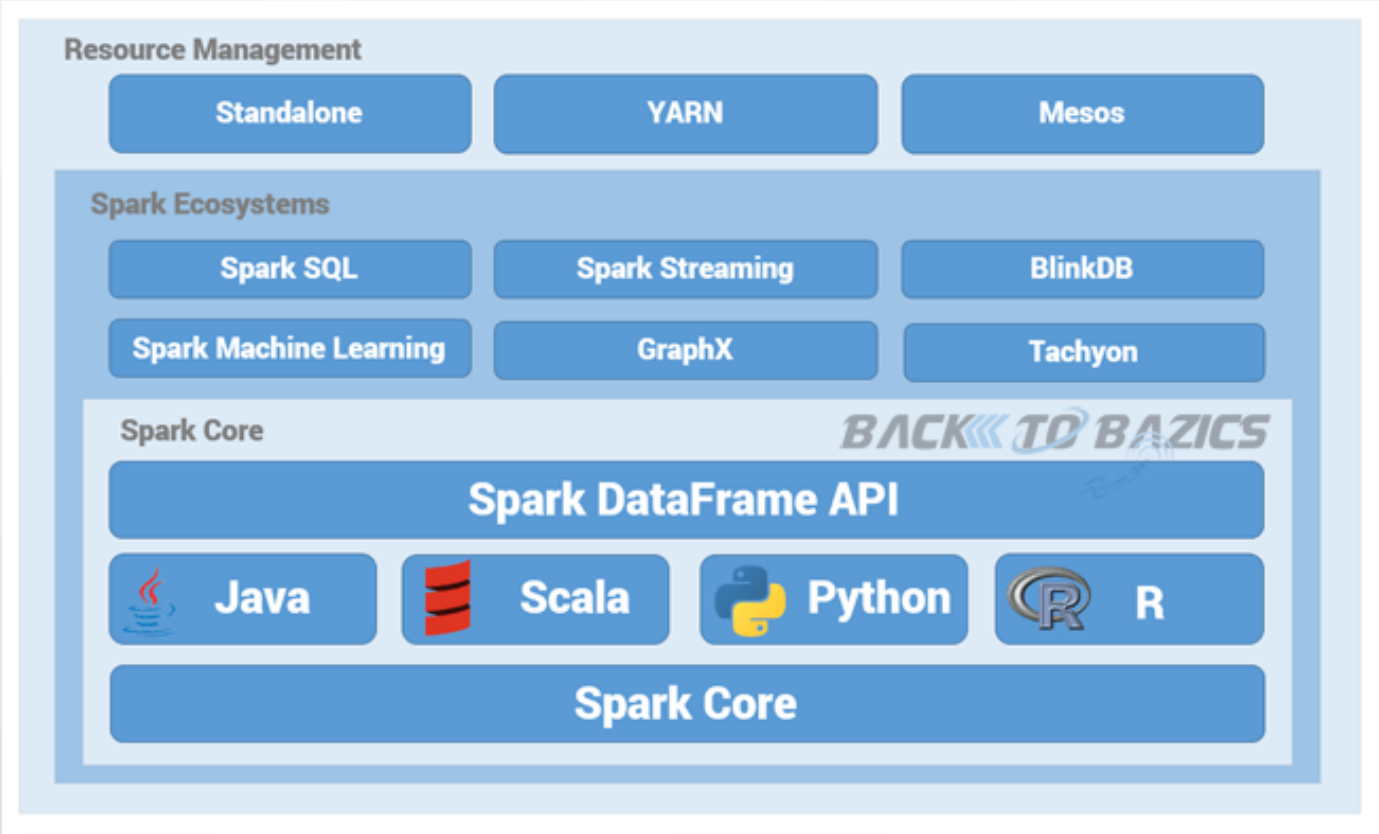
Resource Management
Standalone
YARN
Mesos
Spark Ecosystems
Spark SQL
Spark Streaming
BlinkDB
Spark Machine Learning
GraphX
Tachyon
Spark Core
BACK TO BAZICS
Spark DataFrame API
Java
Scala
Python
R
Spark Core

Spark Streaming

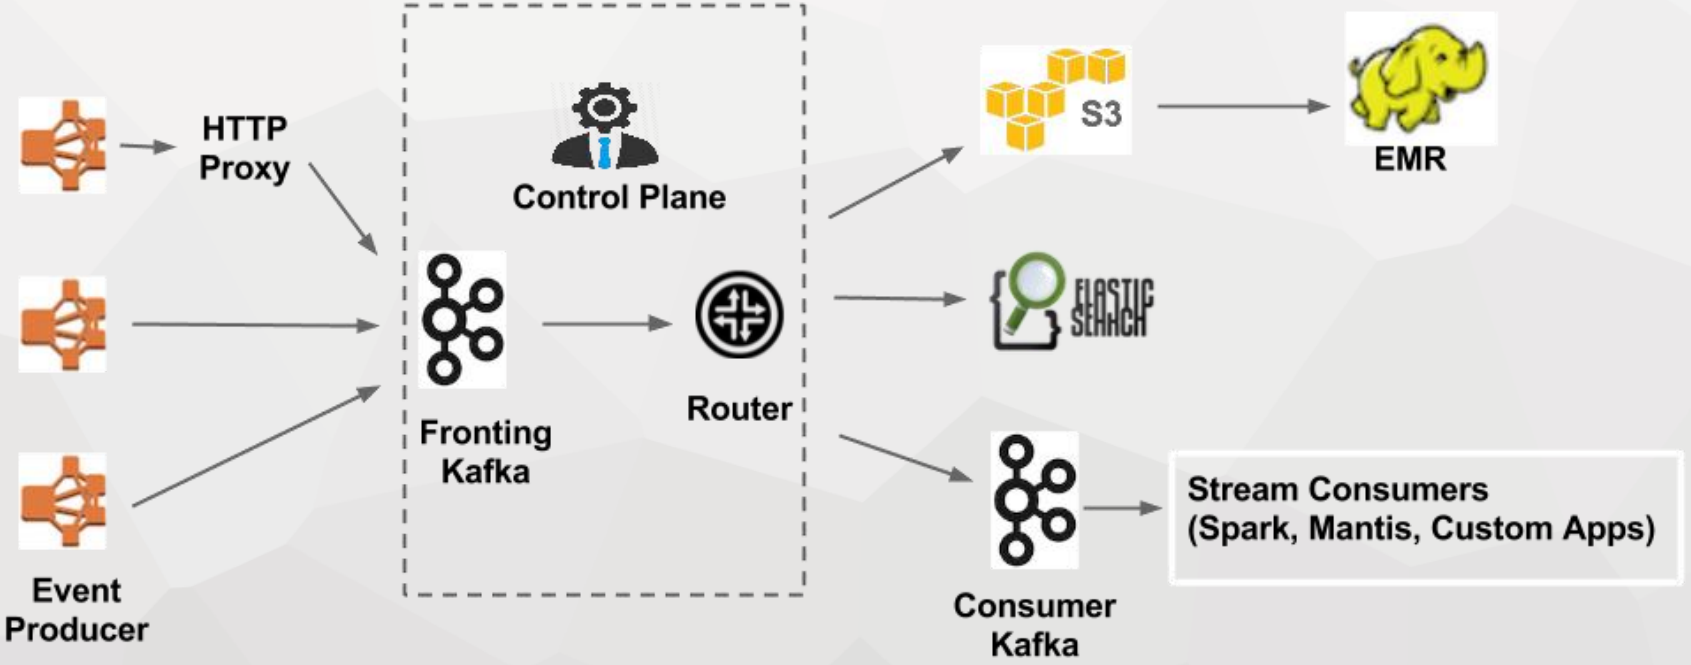
HTTP
Proxy
Control Plane
S3
EMR
Fronting
Kafka
Router
Event
Producer
Consumer
Kafka
Stream Consumers
(Spark, Mantis, Custom Apps)

ELK (ElasticSearch LogStash Kibana)

Search Engine

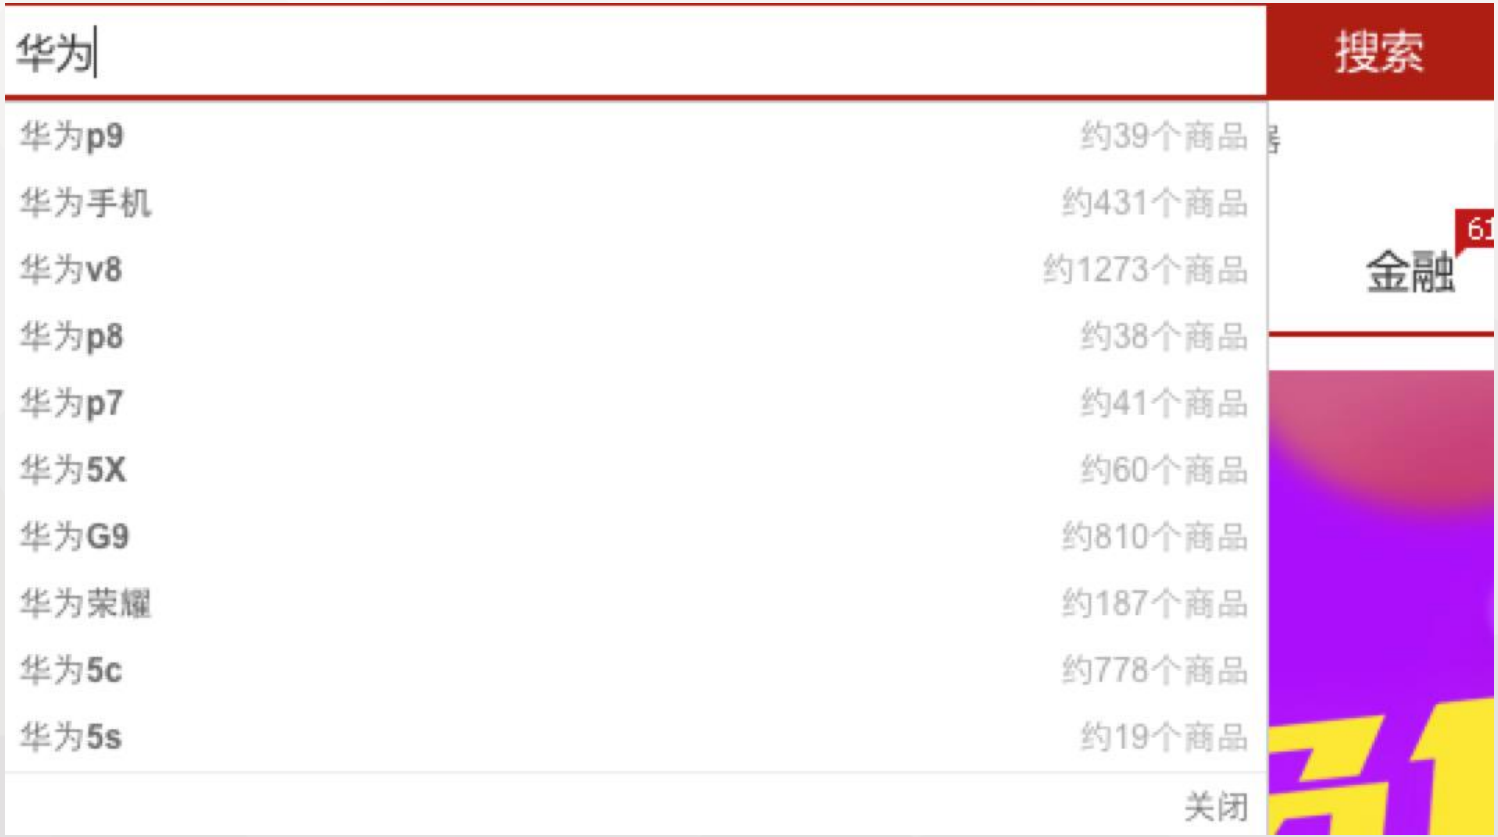

Spark Streaming

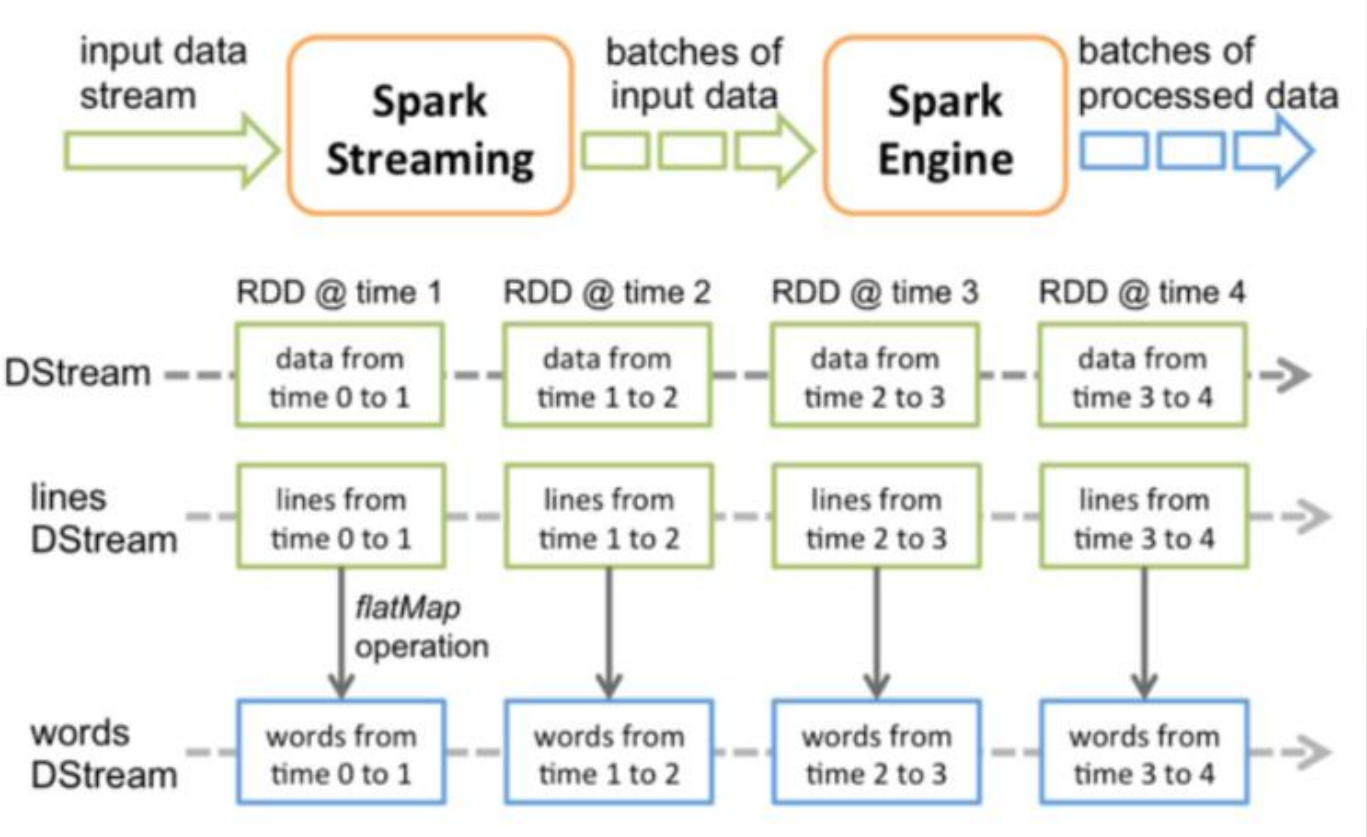

Windows based operations

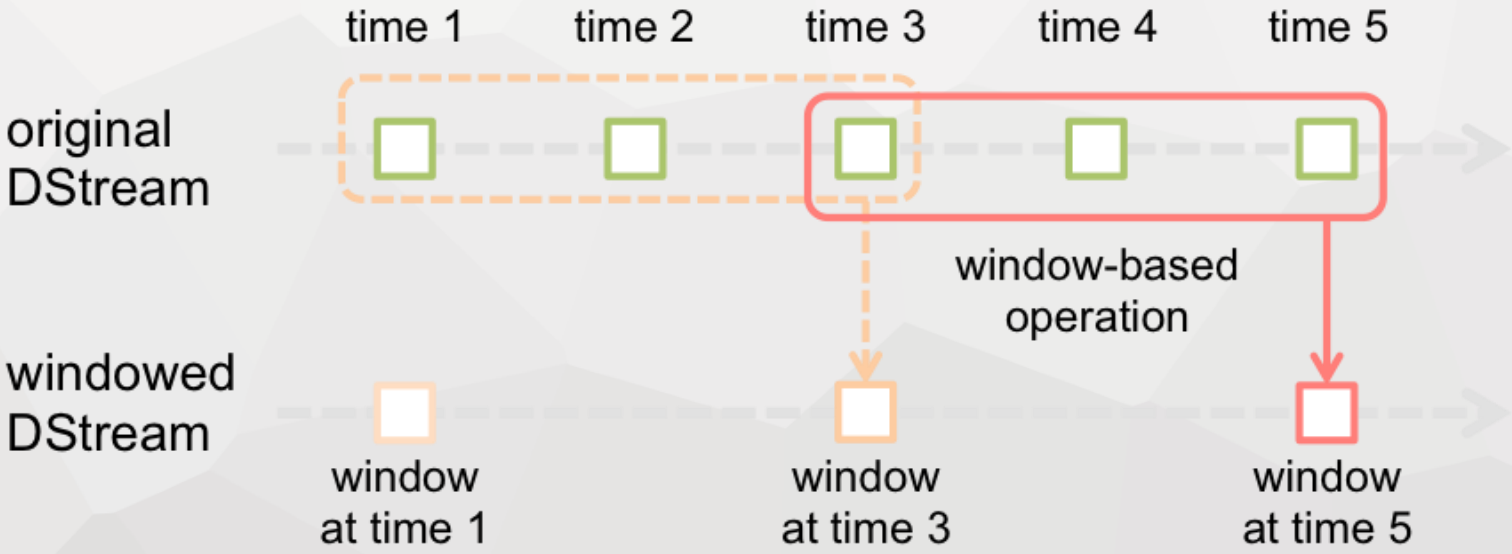

Real-Time Traffic

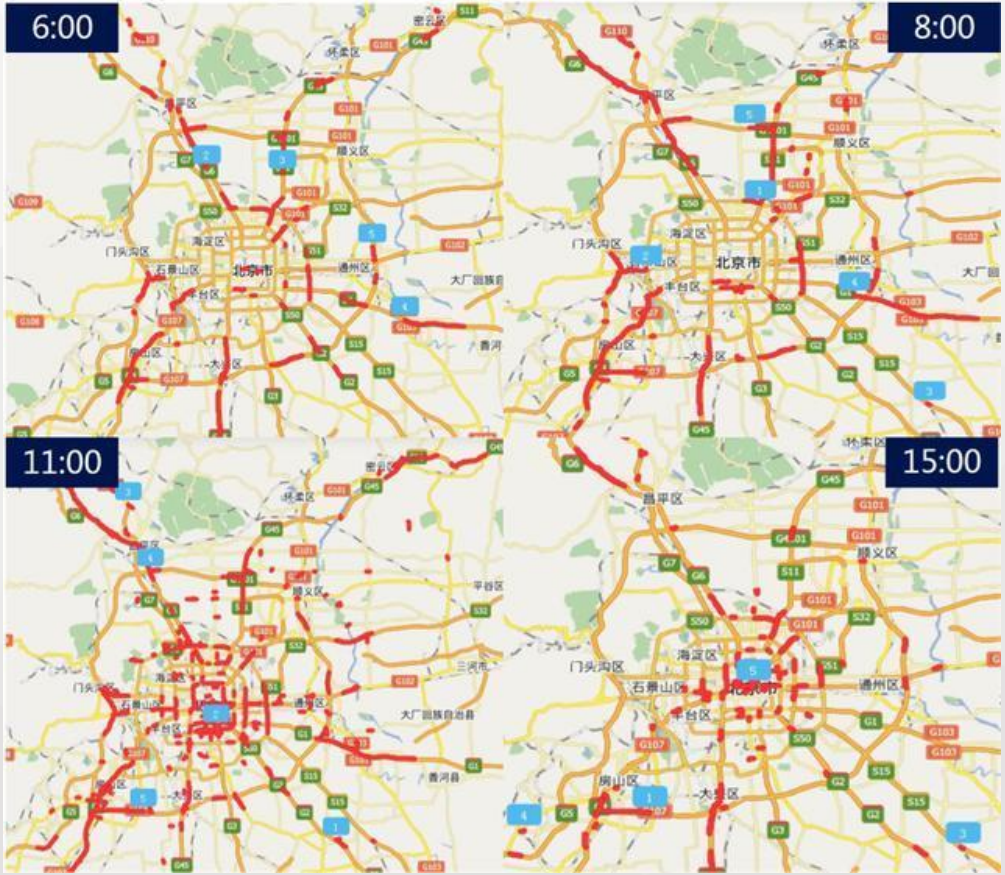

Spark Machine Learning

Ecosystems

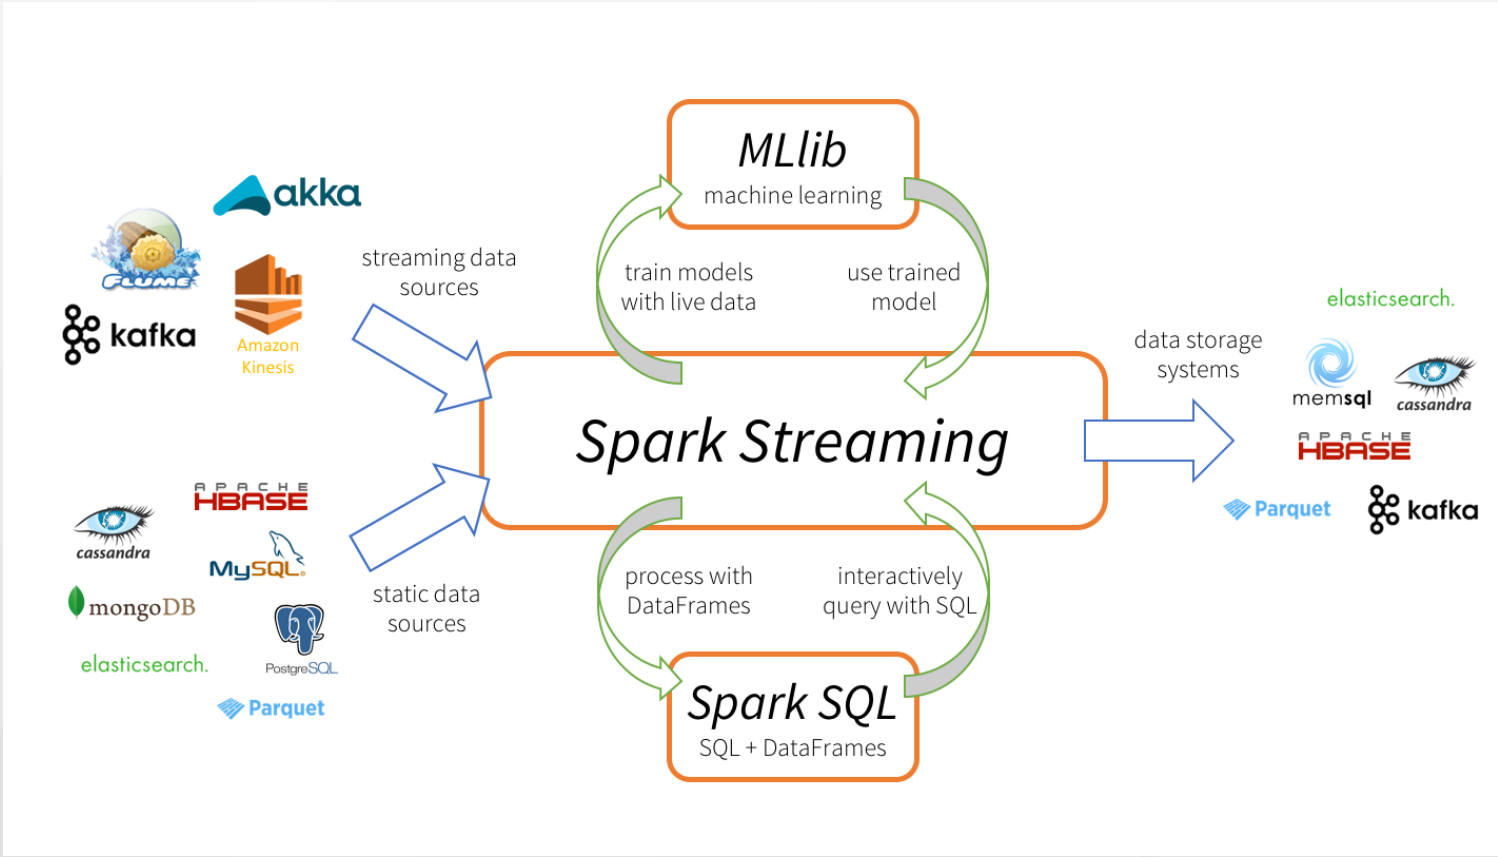
akka
flume
kafka
Amazon Kinesis
streaming data sources
cassandra
APACHE HBASE
MySQL
mongoDB
PostgreSQL
elasticsearch.
Parquet
static data sources
MLlib
machine learning
train models with live data
use trained model
Spark Streaming
process with DataFrames
interactively query with SQL
Spark SQL
SQL + DataFrames
data storage systems
elasticsearch.
memsql
cassandra
APACHE HBASE
Parquet
kafka

Ecosystems

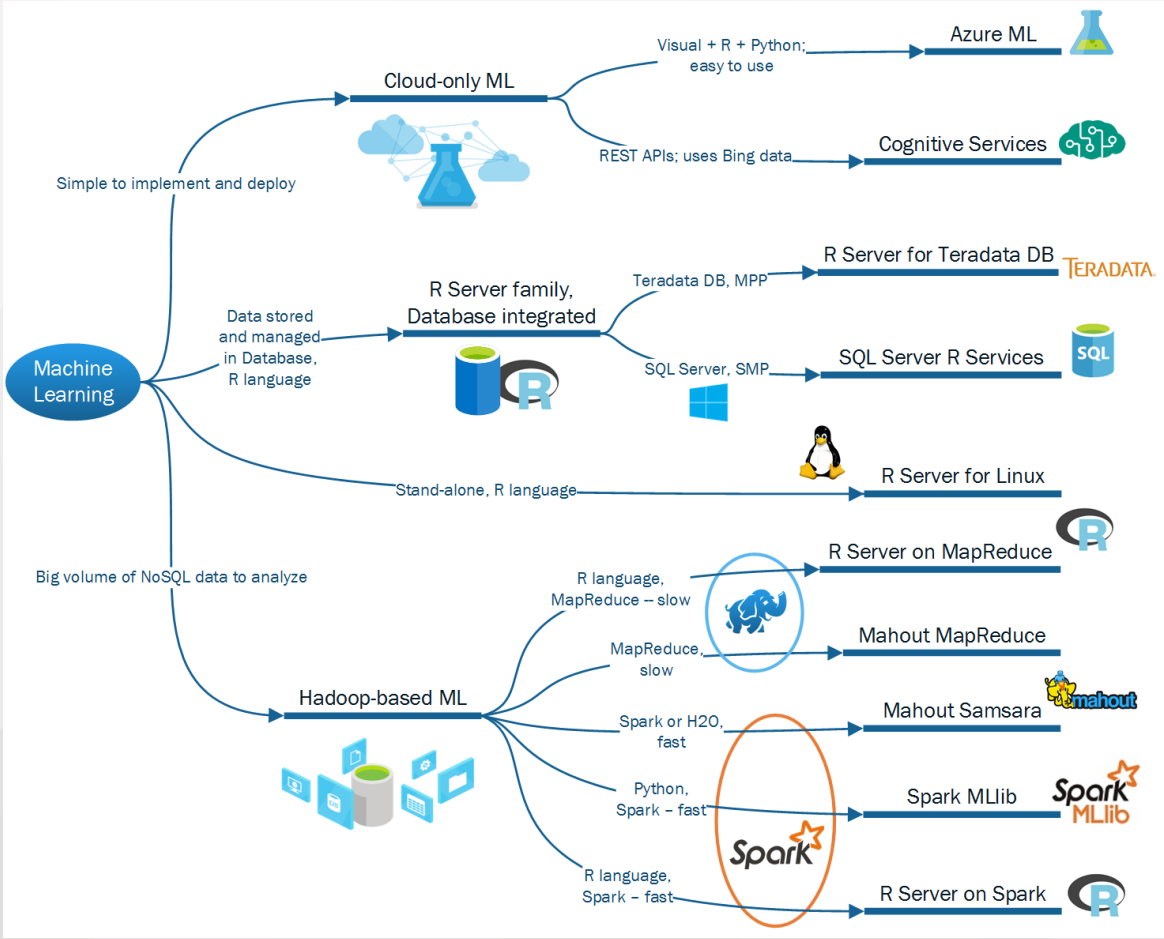
Machine Learning
Simple to implement and deploy
Cloud-only ML
Visual + R + Python; easy to use
Azure ML
REST APIs; uses Bing data
Cognitive Services
Data stored and managed in Database, R language
R Server family, Database integrated
Teradata DB, MPP
R Server for Teradata DB
TERADATA
SQL Server, SMP
SQL Server R Services
SQL
Stand-alone, R language
R Server for Linux
Big volume of NoSQL data to analyze
Hadoop-based ML
R language, MapReduce – slow
R Server on MapReduce
MapReduce, slow
Mahout MapReduce
Spark or H2O, fast
Mahout Samsara
mahout
Python, Spark – fast
Spark MLlib
Spark
R language, Spark – fast
R Server on Spark

Use Case : Credit Approval

Learning Flow for Credit Approval

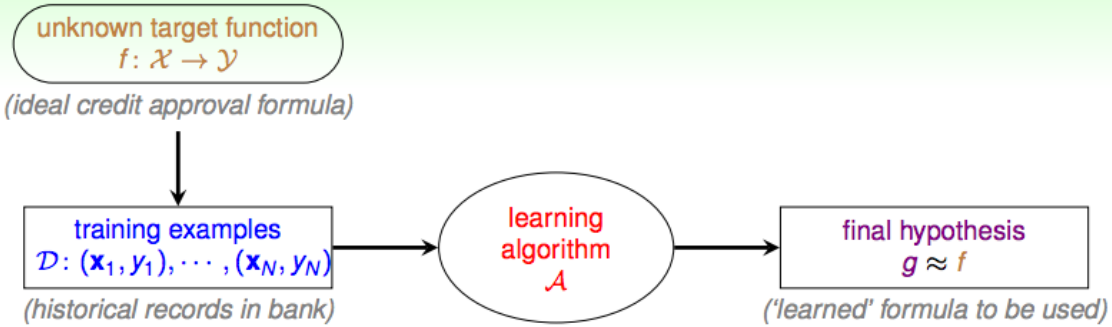

- target f **unknown** (i.e. no programmable definition)
- hypothesis g hopefully $\approx f$ but possibly **different** from f (perfection 'impossible' when f unknown)

What does g look like?

Use Case : WeChat

Use Case : Hatsune Miku Upscale

Spark GraphX

GraphX

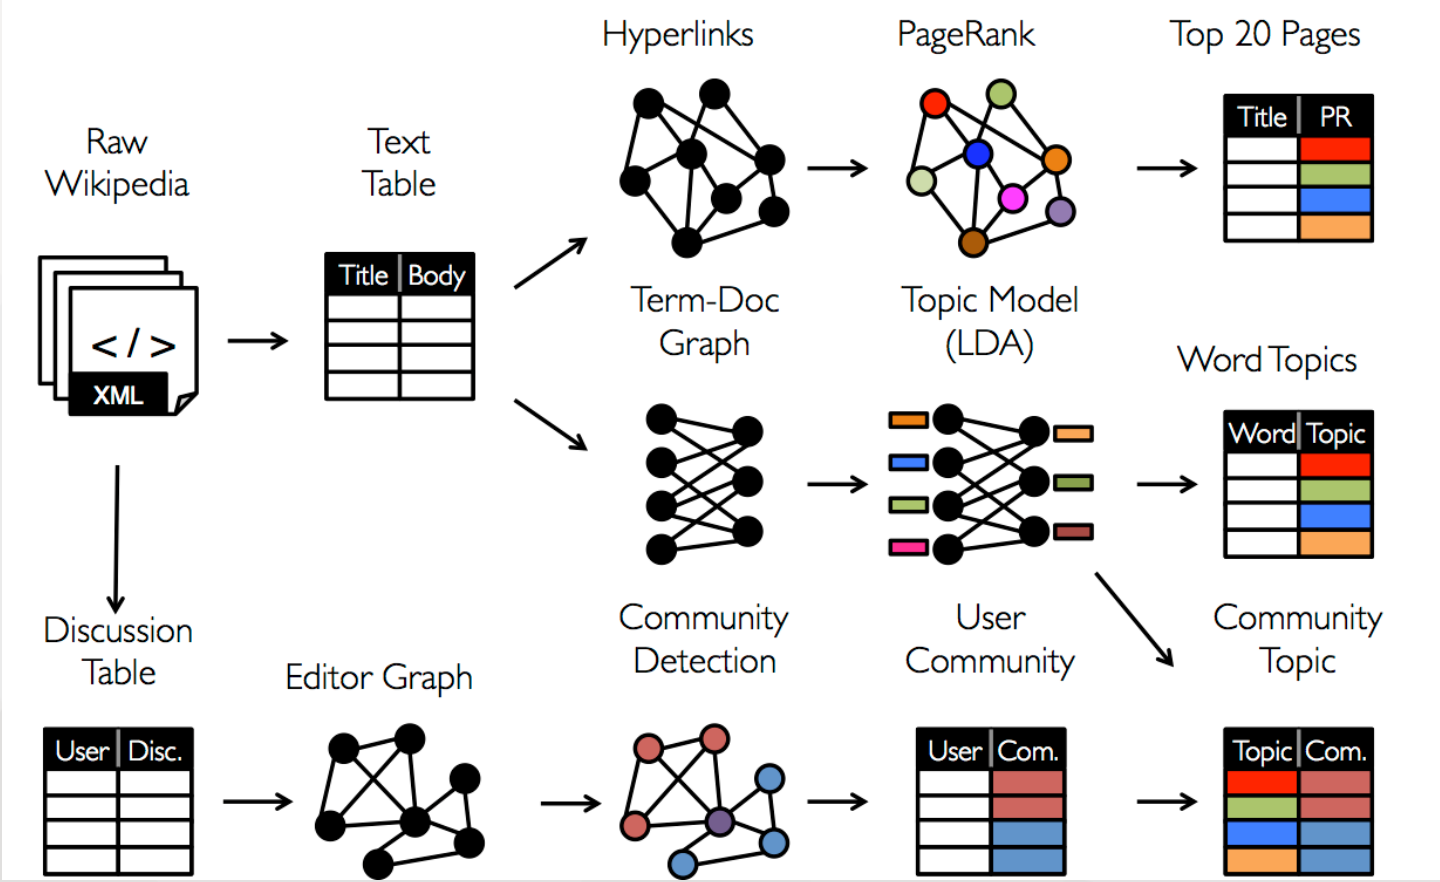
Hyperlinks
PageRank
Top 20 Pages
Raw Wikipedia
Text Table
Title
PR
Title
Body
XML
Term-Doc Graph
Topic Model (LDA)
Word Topics
Word
Topic
Discussion Table
Community Detection
User Community
Community Topic
Editor Graph
User
Disc.
User
Com.
Topic
Com.

Weibo

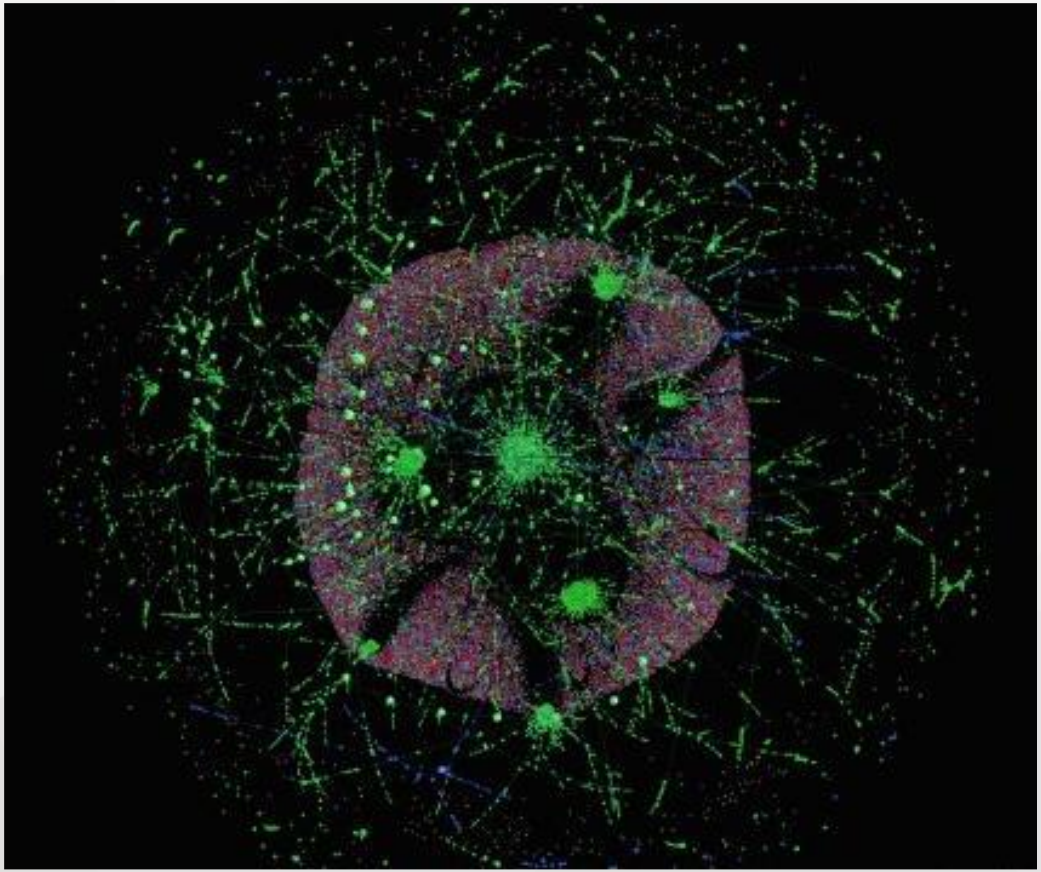

Social Network Site community

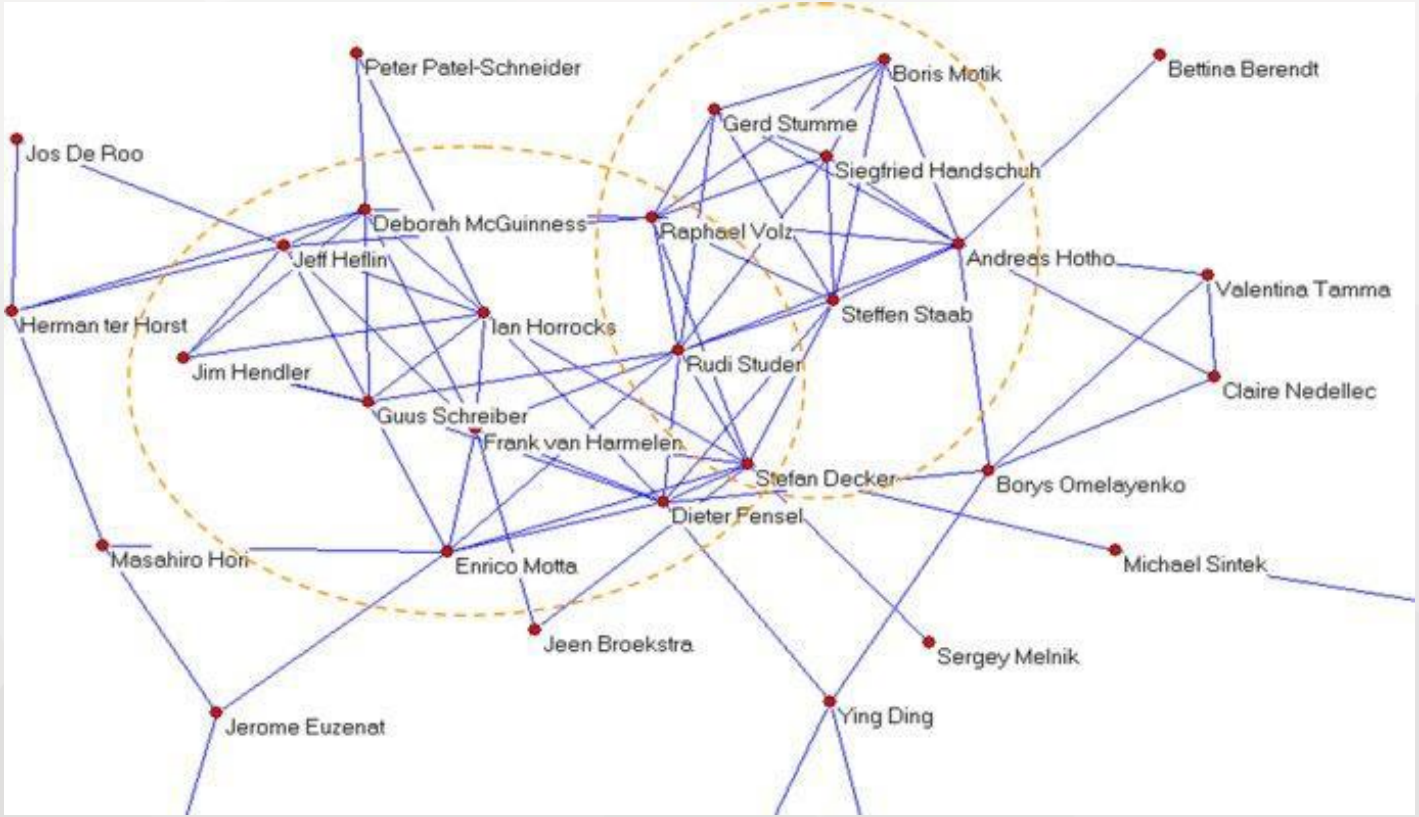


The End